



Computational Discrete Optimization

CO 353



Chaitanya Swamy

Preface

Disclaimer Much of the information on this set of notes is transcribed directly/indirectly from the lectures of CO 353 during Winter 2022 as well as other related resources. I do not make any warranties about the completeness, reliability and accuracy of this set of notes. Use at your own risk.

Discrete optimization problems are underlying decisions that have a discrete flavor, e.g., YES/NO or $\{0,1\}$ decisions.

The focus in this course will be on algorithms, modelling. Broad classes of problems that we will study are network connectivity problems, location problems, general integer programs.

References:

[KT] Algorithm Design, Jon Kleinberg and Eva Tardos, Addison Wesley, 2005.

[V] Approximation Algorithms, Vijay Vazirani, Springer-Verlag, 2001.

[WS] The Design of Approximation Algorithm, David Williamson and David Shmoys, Cambridge University Press, 2011.

For any questions, send me an email via <https://notes.sibeliusp.com/contact>.

You can find my notes for other courses on <https://notes.sibeliusp.com/>.

Sibelius Peng

Contents

Preface	1
1 Graph Algorithms	4
1.1 Definitions, Notations & Terminology	4
1.2 Shortest paths: Dijkstra's algorithm	4
1.3 Running time and Efficient Algorithms	6
2 Graph Algorithms cont'd	8
2.1 Minimum Spanning Trees	8
2.2 Cut property	9
2.3 Prim's Algorithm	9
2.4 Kruskal's algorithm	11
2.5 Application to Clustering	12
3 Graph Algorithms - cont'd	14
3.1 Arborescences	14
3.2 Min Cost Arborescence (MCA) Problem	15
3.3 Edmond's algorithm	16
4 Matroids	21
4.1 Introduction	21
4.2 Max-weight independent set (MWIS) problem	22
4.3 Applications of Matroid Optimization	25
5 Steiner Tree & Computational Complexity	27
5.1 Minimum Steiner Tree Problem	27
5.2 MST-based algorithm	28
5.3 Computational Complexity	30
5.4 Polynomial-Time Reductions	30
6 P & NP	33
6.1 NP	33
6.2 NP-hard and NP-complete	34
7 NP-Hard and NP-Complete	36
8 Approximation Algorithm & Primal Dual Algorithm	39
8.1 Approximation Algorithm Design	39
8.2 Network Connectivity Problems	41
8.3 LP-relaxation for f -connectivity and its dual	41

9	Primal Dual Algorithm cont'd	44
9.1	Primal dual algorithm for f -connectivity	44
9.2	Analysis	47
10	Set Cover & Uncapacitated Facility Location	50
10.1	LP-relaxation for SET COVER and its dual	50
10.2	LP-Rounding Algorithms for Set Cover	52
10.3	Uncapacitated Facility Location	54
11	UFL	55
11.1	Economic Interpretation of Dual LP	55
11.2	Metric UFL	56

Graph Algorithms

1.1 Definitions, Notations & Terminology

A **graph** is a tuple (V, E) , where V is set of **nodes/vertices**, E is set of **edges**, where edges **joins** two nodes.

If e is an edge that joins nodes u, v , then we denote this by $e = uv$. u, v are called **ends** of e . e is **incident** to nodes u, v . We are not allowing parallel edges, i.e., $e = uv$, and $e' = u'v'$ are distinct edges, then $\{u, v\} \neq \{u', v'\}$.

An **u - v path** in $G = (V, E)$ where $u, v \in V, u \neq v$, is a sequence of nodes $u_1 = u, u_2, \dots, u_k, u_{k+1} = v$, where $u_i u_{i+1} \in E \forall i = 1, \dots, k$. A **cycle** in G is a sequence of nodes $u_1, u_2, \dots, u_k, u_{k+1} = u_1$ where $u_i u_{i+1} \in E \forall i = 1, \dots, k$, and u_i 's are distinct. Since there are no parallel edges, we can also identify a path/cycle by its sequence of $u_i u_{i+1}$ edges. So we will often refer to a path/cycle as a set of edges.

A graph G is **connected** if it has a $u - v$ path $\forall u, v \in V (u \neq v)$. G is acyclic if G does not have a cycle. A **tree** is a connected, acyclic graph.

Let $G = (V, E)$ be a connected graph, and $T = (V_T, E_T)$ be a tree. If $E_T \subseteq E$ and $V_T = V$, then we say that T is a **spanning tree** of G .

If C is a cycle, and $e \in C$, then $C - \{e\}$ still connects all nodes of C . So if G is a connected graph, and it contains a cycle C , and $e \in C$, then $G - \{e\} := (V, E - \{e\})$ is a connected graph. Hence, a spanning tree of G is a minimal connected subgraph of G . I.e., if $T = (V, F)$ where $F \subseteq E$ is a minimal set such that (V, F) is connected, then T is a spanning tree of G . If $T = (V, F)$ contains a cycle, then F is not minimal.

In **directed graph**, each edge has a direction, and goes **from** a node **to** another node.

1.2 Shortest paths: Dijkstra's algorithm

Problem Given a directed graph $G = (V, E)$ with edge costs $\{c_e \geq 0\}$ and a node $s \in V$, find the shortest path from s to all other nodes. The "shortest" path means path with the smallest total edge cost under the c_e edge costs.

Notation For a path P , let $c(P) := \sum_{e \in P} c_e$ denote the total cost of P . Let $d(u) = \min_{P: P \text{ a } s \rightarrow u \text{ path}} c(P)$, which is shortest path (SP) distance from s to u . If $u \rightarrow v$ is an edge of G , we have

$$d(v) \leq d(u) + c_{u,v} \quad (\clubsuit)$$

Dijkstra's Algorithm

The idea is to maintain a set of explored vertices, and we want to expand this set. Then we can make use of (\clubsuit) to estimate the shortest path from s to v , a vertex to be added to the set. We will maintain a label $\ell(v)$ for all $v \notin A$, which is our current estimate for the $s \rightarrow v$ shortest path distance.

Given Directed graph $G = (V, E)$, $s \in V$, edge costs $\{c_e \geq 0\}$.

Algorithm 1: Dijkstra's Algorithm

```

1 Initialize  $A \leftarrow \{s\}$ ,  $d(s) = 0$ ,  $\ell(v) \leftarrow \infty \forall v \notin A$ .
2 while  $A \neq V$  do
3   For all  $v \notin A$  such that  $\exists u \in A$  with edge  $u \rightarrow v$ , update
      
$$\ell(v) = \min \left\{ \ell(v), \min_{u \in A: (u,v) \in E} (d(u) + c_{u,v}) \right\}$$

4   Select  $w \in V - A$  such that  $\ell(w)$  has minimum  $\ell(v)$  value among all  $v \notin A$ .
5   Update  $A \leftarrow A \cup \{w\}$ , set  $d(w) = \ell(w)$ .
```

Remark:

Can obtain actual shortest paths by maintaining along with $\ell(w)$, the node $u \in A$ that determines $\ell(w)$ (i.e., $u \in A$ is s.t. $\ell(w) = d(u) + c_{u,w}$). Call u , the “parent” of w , and $u \rightarrow w$ the parent edge of w .

The shortest paths obtained via previous point have a special structure: every node $w \neq s$ has exactly one edge entering it, and there are no cycles, i.e., we have something like “directed” tree. And we denote shortest-path tree: directed tree returned by algorithm.

Also note that $\ell(v)$ in

$$\ell(v) = \min \left\{ \ell(v), \min_{u \in A: (u,v) \in E} (d(u) + c_{u,v}) \right\}$$

is redundant, since

$$\min_{u \in A: (u,v) \in E} (d(u) + c_{u,v})$$

term only decreases as the set A only grows.

Correctness

We may assume that there exists $s \rightarrow u$ path in $G \forall u \in V$. And it's easy to modify Dijkstra's algorithm to detect if this assumption holds, and get shortest path distances from s to all nodes reachable from s .

Let $d^{\text{Alg}}(v)$: d -value computed by algorithm. Recall $d(v)$ is the shortest path distance from s to v . The goal then is to show that for all $v \in V$, $d^{\text{Alg}}(v) = d(v)$. Clearly this is satisfied when $v = s$.

Assume we have correctly computed shortest path distances for all $u \in A$, $\ell(v)$ is the length of the shortest path P such that *last edge of P (which enters v) comes from a node in A* .

Why? Consider such a path P . Let $u \rightarrow v$ be the last edge of P . So $u \in A$, $d^{\text{Alg}}(u) = d(u)$,

$$c(P) \geq d(u) + c_{u,v} = d^{\text{Alg}}(u) + c_{u,v} \geq \ell(v)$$

and last inequality is by the definition of $\ell(v)$.

Theorem 1.1

If w is added to A in line 5 of the algorithm, then $d^{\text{Alg}}(w) = d(w)$. (I.e., we have computed shortest path distance from s to w .)

Proof:

Assume we have correctly computed shortest path distance $\forall u \in A$. Consider an arbitrary $s \rightarrow w$ path P . Let u be the last node on P that lies in A . Let v be the node on P after u (so $v \notin A$). Let P' be the $s \rightarrow v$ portion of P . Then

$$c(P) \geq c(P') \geq d(u) + c_{u,v} = d^{\text{Alg}}(u) + c_{u,v} \geq \ell(v) \geq \ell(w)$$

where the last equality is by the definition of w in the line 4. □

Then following parent edges gives an $s \rightarrow w$ path of length $= \ell(w) = d^{\text{Alg}}(w)$.

1.3 Running time and Efficient Algorithms

The goal in this course is to design efficient algorithms. What does efficient mean? The short answer is “reasonable” running time.

Running time is number of elementary operations performed by algorithm as a function of input size. **Elementary operations** includes basic arithmetic (e.g., addition), comparisons (is $x < y$?), simple logical constructs (i.e., if-then-else), assignments. **Input size** is the number of *bits* needed to specify the input. Note that number of bits need to specify a number $x \geq 0$, x integer is roughly $\log_2 x$, which is much smaller than x itself.

For example, the size of an input of the Dijkstra’s algorithm, $G = (V, E), \{c_e\}_{e \in E}$ is usually taken to be approximately $|V| + |E| + \sum_{e \in E} \log_2 c_e$.

Reasonable running time, i.e., efficient algorithm means that running time that is **polynomial function** of input size. In order to specify running time & input size in a convenient, compact way, we will use $O(\cdot)$ notation.

Given two functions: $f, g : \mathbb{R}_+ \mapsto \mathbb{R}_+$, we say that $f(n) = O(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $f(n) \leq c \cdot g(n)$ for all $n \geq n_0$.

Here are some examples:

$$\begin{aligned} n &= O(n) \\ 2n + 10 &= O(n) \\ 3n &= O(n^2) \\ \alpha n^c + \beta &= O(n^d) \\ n \log_2 n &= O(n^2) \\ \log_2 n &= O(\log_{10} n) \\ 2^n &= O(3^n) \end{aligned}$$

$f(n) = O(1)$ means $f(n) \leq c$ for all $n \geq n_0$. $f(n) = O(n^{O(1)})$ is shorthand for $f(n)$ is bounded by some (fixed) polynomial function of n : $f(n) \leq d \cdot n^c$.

An algorithm with running time $f(n)$, where n is input size, is **efficient** if $f(n)$ is bounded by a polynomial function of n , i.e., $f(n) = O(n^{O(1)})$.

Now we can examine the running time of Dijkstra's algorithm (removing unnecessary $\ell(v)$ in line 3). Let $m = |E|$, $n = |V|$. We observe that there are n iterations of while loop. In each iteration:

1. Computing $\ell(v)$ takes $O(d^{\text{in}}(v))$ time where $d^{\text{in}}(v)$ is the number of edges entering v .
2. Computing $\ell(v) \forall v$ takes $O(m)$ time since $\sum_{v \in V} d^{\text{in}}(v) = m$.
3. Line 4 takes $O(n)$ time.
4. Line 5 takes $O(1)$ time.

Each iteration takes $O(m + n)$ time. This is $O(m)$ if we assume there exists an $s \rightarrow v$ path $\forall v \in V$ since then $m \geq n - 1$, so $n = O(m)$. Then the running time of algorithm is $O(mn)$ which is a polynomial function of input size.

However, we can have a better implementation. Observe that if $\{u \in A : (u, v \in E)\}$ does not change across iterations, then $\ell(v)$ does not change. So instead of recomputing $\ell(v)$ for all $v \notin A$, we do the following:

When we pick $w \notin A$ to add to A , we only update $\ell(v)$ for all $v \notin A$ such that $(w, v) \in E$, and set $\ell^{\text{new}}(v) = \min(\ell^{\text{old}}(v), d(w) + c_{w,v})$.

So the steps inside of while loop change as: [Let w^* be the last node added to A . Initially $w^* = s$.]

- (a) For every edge (w^*, v) , where $v \notin A$, update

$$\ell(v) = \min(\ell(v), d(w^*) + c_{w^*,v})$$

and we call this DecreaseKey operation.

- (b) Find $w \notin A$ with minimum $\ell(\cdot)$ value. We call this ExtractMin operation.

- (c) Update $A \leftarrow A \cup \{w\}$, $d(w) = \ell(w)$, $w^* = w$.

Across all iterations, we examine each edge (u, v) at most once in step (a) above (in the iteration when $w^* = u, v \notin A$). So across all iterations, $\leq m$ DecreaseKey operations, $\leq n$ ExtractMin operations.

Then we can use a simple array to store $\ell(\cdot)$ values. Note that DecreaseKey is $O(1)$, and ExtractMin operation is $O(n)$. Thus the running time = $O(m + n^2) = O(n^2)$.

There exist data structures such as priority queue, under which DecreaseKey and ExtractMin take $O(\log n)$. Then the running time is then $O(m \log n)$.

There exists a data structure called Fibonacci heaps, under which DecreaseKey is $O(1)$, and ExtractMin operation is $O(\log n)$. Then the running time is $O(m + n \log n)$.

Graph Algorithms cont'd

2.1 Minimum Spanning Trees

MST Problem Given a connected, undirected graph $G = (V, E)$, edge costs $\{c_e\}_{e \in E}$. Find a spanning tree of G of minimum total edge cost.

We say “ T is a spanning tree” is equivalent to “ T is the edge set of a spanning tree”. We denote the cost of T by $c(T) := \sum_{e \in T} c_e$.

Note:

The c_e 's could be positive, zero, or negative.

If all c_e 's are ≥ 0 , then can equivalently define the MST problem as: find the min-cost connected spanning subgraph of G . Because there is always an optimal solution that is minimal connected spanning subgraph of G .

Theorem 2.1: Fundamental Theorem about trees

Let $T = (V, F)$ be a graph, and $n = |V|$. The following are equivalent:

- (a) T is a tree (i.e., connected, acyclic)
- (b) T is connected, has $n - 1$ edges.
- (c) T is acyclic, has $n - 1$ edges.

Proof:

(a) \Rightarrow (b) Pick some $r \in V$ as root node. Root T at r , i.e., draw T as hanging off of r . For each $v \neq r$, there is a unique edge uv of T (incident to v) such that u is closer to r than v , and we call uv the parent edge of v .

These parent edges cover T , and number of parent edges $= n - 1$, since each $v \neq r$ has a unique parent edge.

(b) \Rightarrow (a) T is connected. Let T' be a spanning tree of T . So by (a) \Rightarrow (b), we know that T' has $n - 1$ edges. But T has $n - 1$ edges. So $T = T'$, so T is a tree. \square

2.2 Cut property

Notation Let $v \in V$. $\delta(v)$ denotes the set of edges incident to v . Let $S \subseteq V$, $\delta(S) := \{uv \in E : u \in S, v \notin S\}$. In other words, $\delta(S)$ denotes the “boundary” of S .

Now we assume that all edges costs are distinct. Fix some node $s \in V$. Let $e \in \delta(S)$ have the smallest edge cost among edges in $\delta(S)$. Is e in some MST? In fact, e is in every MST.

cut

A cut is any partition $(A, V - A)$ of the vertex set V , where $A \neq \emptyset, A \subsetneq V$.


Edges crossing the cut are edges in $\delta(A)$ ($= \delta(V - A)$). If $F \subseteq E$, we say F crosses the cut to mean $F \cap \delta(A) \neq \emptyset$.

Lemma 2.2: Cut property

Consider any cut (A, B) , where $B = V - A$. If e is the (unique) min-cost edge across the cut, then e belongs to every MST.

Proof (via an exchange argument):

Suppose T is an MST such that $e \notin T$. We will show that we can find another spanning tree T' (that contains e) such that $c(T') < c(T)$, then a contradiction.



$T \cup \{e\}$ contains a cycle C that contains e . And this is because $T \cup \{e\}$ is connected and has n edges, then it can't be acyclic. Here is a *basic fact*: if a cycle crosses a cut, it crosses the cut at least twice. So $\exists e' \in C \cap \delta(A)$, $e' \neq e$. By definition of e , $c_e < c_{e'}$. And $e' \in T$.

Consider $T' = T \cup \{e\} \setminus \{e'\}$. We claim that T' is a spanning tree. T' is connected since $e' \in$ cycle in $T \cup \{e\}$ and T' has $n - 1$ edges. Then we have

$$c(T') = c(T) + c_e - c_{e'} < c(T)$$

a contradiction. □

2.3 Prim's Algorithm


Now we can use cut property as the basis of the greedy algorithm.

Algorithm 2: Prim's Algorithm

- 1 Pick an arbitrary “seed” node $s \in V$.
- 2 Initialize $A \leftarrow \{s\}$, $T \leftarrow \emptyset$.
- 3 **while** $A \neq V$ **do**
- 4 Choose $e = uv \in \delta(A)$ with smallest cost, where $u \in A, v \notin A$.
- 5 $A \leftarrow A \cup \{v\}$, $T \leftarrow T \cup \{e\}$.
- 6 **return** T

Example:

Blue lines are the output of Prim's algorithm.

**Theorem 2.3**

Prim's algorithm correctly computes an MST.

Proof:

Let T be the edge-set returned by Prim's algorithm.

- T is a spanning tree. T is connected, since every node is connected to s in T . Also, T has $n - 1$ edges. Thus T is a spanning tree.
- T is a MST. Every $e \in T$ belongs to every MST by the cut property, since it is the min-cost edge across some cut. So $T \subseteq$ every MST. But T is itself a spanning tree, so T is MST.

□

Corollary

T is the unique MST.

If edge costs are not distinct, then Prim's algorithm still returns an MST; there could be multiple MSTs.

Implementation & Running Time

Implementation will be similar to Dijkstra's algorithm. For every unexplored node $v \notin A$, maintain a "key" $a(v) = \min_{e=uv: u \in A} c_e$. So in each iteration, we choose $w \notin A$ with smallest $a(\cdot)$ value similar to Dijkstra, let w^* be the last node added to A . In each iteration

- For each edge w^*v , where $v \notin A$, update $a(v) = \min\{a(v), c_{w^*v}\}$. **DecKey**
- Find $w \in V - A$ with smallest $a(\cdot)$ value. **ExtractMin**
- Set $A \leftarrow A \cup \{w\}$, and $T \leftarrow T \cup \{uw\}$, where $u \in A$, and $c_{uw} = a(w)$.

As in Dijkstra's algorithm, across all iterations:

- n **ExtractMin** operations
- m **DecKey** operations

So the running time is

- $O(m + n^2)$ using a simple array to store keys (**DecKey** $O(1)$, **ExtractMin** $O(n)$)
- $O(m + n \log n)$ using a sophisticated data structure like Fibonacci Heaps (**DecKey** $O(1)$, **ExtractMin** $O(\log n)$)

2.4 Kruskal's algorithm

Kruskal's algorithm finds a MST. In some level, it is more greedy and intuitive than Prim's algorithm. The idea is to keep the edge costs in increasing order, and add edges to the set one by one, as long as no cycle are introduced.

Algorithm 3: Kruskal's algorithm

```

1 Sort the edges in increasing order of cost.
2 Initialize  $T \leftarrow \emptyset$ .
3 for each edge  $e$  in sorted order do
4   if  $T \cup \{e\}$  does not have a cycle then
5      $T \leftarrow T \cup \{e\}$ 
6 return  $T$ 

```

Theorem 2.4

Kruskal's algorithm returns the unique MST when all edges costs are distinct.

Proof:


Let T be the edge-set returned by Kruskal. Then, T is acyclic: by construction. Consider a basic fact: a graph is connected $H = (V_H, E_H)$ if and only if $\delta_H(A) \neq \emptyset \forall A : \emptyset \neq A \subsetneq V_H$.

Suppose (V, T) is not connected. Then from the basic fact, $\exists A, \emptyset \neq A \subsetneq V$ such that $\delta(A) \cap T = \emptyset$. But G is connected, so \exists some edge $e \in \delta(A)$. Then $T \cup \{e\}$ is acyclic. So consider the point when Kruskal considers edge e . Let $F \subseteq T$ be set of edges Kruskal has added until then. Then $F \cup \{e\}$ is acyclic, so Kruskal should have added e . Then $e \in T$, a contradiction. So T is a spanning tree.

Consider any edge $e = uv \in T$. Let

$$A = \{w \in V : w \text{ is connected to } u \text{ in } T \text{ at the point when } e \text{ is considered by Kruskal}\}$$

Graph $H = (V, T)$
at the point when
Kruskal considers e



We claim that e is the min-cost edge in $\delta(A)$. Observe that e is the first edge of $\delta(A)$ considered by Kruskal. Hence e is the min-cost edge in $\delta(A)$. By claim, $e \in$ every MST. So $T \subseteq$ every MST. But T is itself a spanning tree. So T is MST. \square

Remark:

We can stop Kruskal when $|T| = |V| - 1$.

Running time Sorting m edges takes $O(m \log m)$. To check if $e = uv$ can be added, we need to check if u, v are in different components of (V, T) at that point. There exist data structures (e.g., Union-Find) for maintaining connected components that allow one to do this $O(\log n)$ time. So total time for step 3 is $O(m \log n)$. Since $m \leq n^2$, the total time for the algorithm is $O(m \log m + m \log n) = O(m \log n)$.

2.5 Application to Clustering

Clustering Given a set of objects, and some notion of similarity/dissimilarity between these objects, divide the objects into groups (called clusters) so that

1. Objects in the same group are “similar” to each other.
2. Objects in different groups are “dissimilar” to each other.

Maximum-Spanning Clustering Given a set $V = \{p_1, \dots, p_n\}$ of objects/points, and pairwise distances $d(p_i, p_j) = d(p_j, p_i) \geq 0 \forall i, j \in [n]$. The goal is to partition V into k clusters C_1, \dots, C_k ($C_i \cap C_j = \emptyset \forall i \neq j$, $\bigcup_{i=1}^k C_i = V$). So as to maximize the minimum inter-cluster spacing, which is equivalent to minimum distance between a pair of points in different clusters. I.e., maximize

$$\min_{\substack{i, j \in [k] \\ i \neq j}} \min_{\substack{p \in C_i \\ q \in C_j}} d(p, q)$$

Algorithm 4: Single-Linkage Clustering

- 1 Start with every point in a separate cluster.
 - 2 Repeatedly merge the 2 clusters with smallest inter-cluster distance^a, until we have k clusters.
-

^aDistance between $C_i, C_j = \min_{p \in C_i, q \in C_j} d(p, q)$


This algorithm is called Single-Linkage Clustering, which is an example of an agglomerative algorithm (i.e., based on merging clusters). Consider a graph G with V as vertex set, and an edge between every pair $p, q \in V$, $p \neq q$, with cost $d(p, q)$.

Note:

Single-Linkage Clustering is exactly Kruskal (merging 2 clusters C_i, C_j due to points $p \in C_i, q \in C_j$) when adding edge pq . *Except* that we stop when there are k components. And this is the same as taking an MST and deleting the $k - 1$ most costly edges, i.e., the $k - 1$ edges that Kruskal could have added last.

Thus equivalently, Run Kruskal (on complete graph with vertex set V , $d(p, q)$ edge costs) but stop when k components remain, which is equivalent to take MST and delete $k - 1$ most costly edges.


Now let's prove the correctness of Single-Linkage Clustering. Let C_1, \dots, C_k be clustering produced by MST - $\{k - 1$ most costly edges $\}$. Let d be the spacing of this clustering.



Observe that

$$\begin{aligned} d &= \text{cost of edge Kruskal would have added next} \\ &= (k - 1)\text{th most costly edge of MST} \end{aligned}$$

If $C = \{C_1, \dots, C_k\}$ is not the optimum, let $C^* = \{C_1^*, \dots, C_k^*\}$ be the optimum clustering. There might exist the case that $C_i \subseteq C_j^*$. As both C and C^* have k partitions, it's not possible to have \subseteq for all k partitions. Since $C \neq C^*$, there is some cluster C_i that intersects at least two clusters of C^* .



Red edges inside C_i all have cost $\leq d$ since these are already added by Kruskal.

So there exists points p, q such that $p, q \in C_i$, but p, q lie in different clusters of C^* . Suppose $p \in C_\ell^*$, $q \in C_r^*$, $\ell \neq r$. Then considering $p - q$ path in C_i , there must be two consecutive nodes u_j, u_{j+1} such that $u_j \in C_\ell^*, u_{j+1} \notin C_\ell^*$. But then spacing of $C^* \leq d(u_j, u_{j+1}) \leq d$ since u_j, u_{j+1} are in different clusters of C^* . So this gives a contradiction since we assumed that $\{C_1, \dots, C_k\}$ is not an optimal clustering, and optimal clustering has spacing strictly larger than d .

Graph Algorithms - cont'd

3.1 Arborescences


Arborescences is directed spanning tree. $G = (V, E)$ a directed graph, and let $r \in V$ be a “root” node.

arborescence rooted at r

An arborescence rooted at r (or rooted out of r) is a subgraph $T = (V, F)$ (so $F \subseteq E$) such that

- there is an $r \rightarrow v$ path in $T \forall v \in V$
- T is a spanning tree if we ignore the directions of the edges. I.e., “undirected version of T ” is a spanning tree.

Example:




Lemma 3.1: Useful alternate characterization of arborescence

$T = (V, F)$ is an arborescence rooted at r if and only if T has no (directed) cycles and every $v \neq r$ has exactly one incoming edge, (and r has no incoming edges).

Proof:

(\Rightarrow) Suppose T is arborescence. Since undirected version of T is a spanning tree, T has no directed cycles, and there exists $r \rightarrow v$ path in T for all $v \in V$ by the definition of arborescence. The path must be unique, otherwise we might encounter the situation like



then we will have a cycle in the undirected version of T . Hence every $v \neq r$ has exactly one incoming edge.

(\Leftarrow) Suppose T has no directed cycles, and every node $v \neq r$ has exactly one incoming edge. For any node $v \neq r$, we construct an $r \rightarrow v$ path as follows: we take v 's unique incoming edge, say $u_1 \rightarrow v$, then u_1 's unique incoming edge, and so on... Since T has no cycle, this process must stop because we have reached r , then we found an $r \rightarrow v$ path.

This also shows that undirected version of T is connected. Also, note that r has no incoming edges in T : suppose T had an edge $v \rightarrow r$, but then $r \rightarrow v$ path in T would create a directed cycle.

So we have shown that (since every $v \neq r$ has exactly one incoming edge) T has $n - 1$ edges. So undirected version of T is connected, has $n - 1$ edges, thus is a spanning tree. \square


3.2 Min Cost Arborescence (MCA) Problem

Given a directed graph $G = (V, E)$, and a root node $r \in V$, and edge costs $\{c_e\}_{e \in E}$, find an arborescence rooted at r of minimum total edge cost.

Assume there exists $r \rightarrow v$ path in G for all $v \in V$, then there exists an arborescence rooted at r .

There are two greedy strategies (inspired by MST)

Strategy 1 Pick the cheapest edge entering a node v (for some fixed node v). Does this edge belong to an MCA? Consider the following example for v . The cheapest edge entering v is of cost 1. However, if we pick this edge, the only arborescence containing edge uv is of cost 106, while MCA is of cost 21.



Strategy 2 Consider some cycle, and delete the most costly edge e of the cycle. Is this valid, i.e., is there an MCA not containing e ? In the above example, we take the cycle $C = u \rightarrow v \rightarrow w \rightarrow u$. If we

delete the most cost edge vw , it's impossible to get MCA on the right.

Let us examine greedy strategy (1) again. Pick cheapest edge entering each node $v \neq r$. Let F^* be set of edges picked. Now let's make two observations.

Observation 1 If (V, F^*) is an arborescence, then it is an MCA. Every arborescence must pick an incoming edge for every $v \neq r$, and F^* is the cheapest way of picking these edges. If (V, F^*) is NOT an arborescence, then by Lemma 3.1, we know that (V, F^*) contains cycle Z , not containing root r .

Observation 2 Suppose for every node $v \neq r$, for each edge e entering v , we subtract a common amount P_v from its cost, i.e., $\forall v \neq r, \forall$ edges e entering v , we set $c'_e = c_e - P_v$. Then for any arborescence T , its c -cost and c' -cost differ by a constant; more precisely, $c(T) - c'(T) = \sum_{v \neq r} P_v$ (simply because T has exactly one edge entering each $v \neq r$ by lemma 3.1).

Corollary

Let $y_v = \min_{(u,v) \in E} c_{u,v}$. For all $v \neq r$, for all edges e entering v , let $c'_e = c_e - y_v$ ^a. Then T is an MCA with respect to $\{c_e\}$ if and only if T is an MCA with respect to $\{c'_e\}$ costs.

^a $c'_e \geq 0$ for all e .

3.3 Edmond's algorithm

Now based on the corollary, imagine we have a cycle $Z \subseteq F^*$. And by observation 2, all edges of F^* (hence Z) have $c'_e = 0$. Then as long as we can reach one node of the cycle, we can take appropriate edges of the cycle to reach all nodes in the cycle without increasing the c' -cost. As we don't care about the inside of the cycle, we can contract the cycle and run the algorithm recursively.


Algorithm 5: Edmond's Algorithm for MCA

Input: Directed graph $G = (V, E)$, root $r \in V$, $\{c_e\}$ edge costs; assume there exists $r \rightarrow v$ path for all $v \in V$


- 0 **if** G has only one node **then return** \emptyset
 - 1 **foreach** $v \neq r$ **do**
 - $y_v := \min_{(u,v) \in E} c_{u,v}$
 - $c'_e := c_e - y_v$ for all e entering v
 - 2 For each $v \neq r$, choose a 0 c' -cost edge entering v . Let F^* be the resulting set of edges.
 - 3 **if** F^* is an arborescence **then return** F^*
 - 4 Otherwise F^* contains a directed cycle. Find a cycle $Z \subseteq F^*$ (not containing r). Contract Z into a single supernode^a to get a graph $G' = (V', E')$.
 - 5 Recursively find an MCA (V', F') in G' with respect to $\{c'_e\}$ edge costs.
 - 6 Extend (V', F') to an arborescence (V, F) in G :
 - Let $v \in Z$ be the node that has an incoming edge in F' .
 - Set $F \leftarrow F' \cup Z \setminus \{\text{edge of } Z \text{ entering } v\}$
 - 7 **return** F
-

^aRemove self-loops, but retain parallel edges


Now let's consider an example.




We then change the cost to c'_e .




Then we mark the zero costs edges to produce F^* . And we can see there are cycles in this set of edges.




We first pick cycle Z_1 and contract it. And do the cost change again.




And we found cycle Z_2 , and contract it, and do the cost change.



Now red edges are arborescence. Recursive calls stop. Then we expand it back.



Then expand again.



which is the output of the algorithm.

Running time

Line 0 takes $O(1)$. Line 1-2 are operations on edges, take $O(m)$. Line 3 on checking arborescence takes $O(n)$: follow incoming edges in F^* to see if we can reach r , either rv path or a cycle. Line 4 takes $O(|Z|) = O(n)$. Line 6 takes $O(|Z|) = O(n)$.

Algorithm makes $O(m)$ elementary operations and a recursive call to a smaller graph; at most n recursive calls. Thus polynomial running time.


Correctness

Lemma 3.2

Suppose we have $\{d_e\}$ edge costs, and a 0 d -cost cycle Z such that $r \notin Z$. Then there exists an MCA (with respect to d -costs) that has exactly one edge entering Z .


Proof:

Let T be any arborescence, where it might have multiple edges entering Z . We will show that T can be modified to an arborescence T' with the stated property and such that $d(T') \leq d(T)$ where $d(T) := \sum_{e \in T} d_e$. T might look like the graph below.



Among all edges $(u', v') \in T$ that enter Z , let (u, v) be such that the $r \rightarrow u$ path in T has the fewest number of edges. Note that the $r \rightarrow u$ path in T has no nodes from Z . Let


$$T' \leftarrow T - \{(u', v') \in T : v' \in Z, u' \notin Z, v' \neq v\} \cup \{\text{edges of } Z \text{ except for the edge entering } v\}$$



Note that $d(T') \leq d(T)$ because all edges we are adding in Z have zero costs. We now have to show T' is an arborescence.

T' has $n - 1$ edges since every $w \neq r$ has exactly one incoming edge. Now it's remained to show that there exists an $r \rightarrow w$ path $T' \forall w \neq r$. Let P be the $r \rightarrow w$ path in T . If P does not contain any node of Z , then P is also an $r \rightarrow w$ path in T' .

Now suppose P contains a node of Z . Let w' be the last node of P in Z . Then T' contains:



- $r \rightarrow v$ path in T
- $v \rightarrow w'$ portion of Z
- $w' \rightarrow w$ portion of P

Concatenating these gives an $r \rightarrow w$ path in T' .

This also shows that undirected version of T' is connected, and has $n - 1$ edges. Thus the undirected version is a spanning tree. \square

Observation Any arborescence in G that enters cycle Z in step 4 of algorithm exactly once, yields an arborescence in G' and VICE VERSA, and these two arborescence have the same c' -cost.

Theorem 3.3

Edmond's algorithm finds an MCA of G .

Proof:

By induction on $|V|$. Base cases where $|V| = 1$ or 2 are clearly true. Suppose inductively, algorithm finds an MCA T' in G' with respect to c' -edge costs.

Let T^* be an MCA in G with respect to c' -edge costs that enters Z exactly once, which exists by lemma 3.2. Let $T^{*'}$ be arborescence obtained from T^* for graph G' .

We have

$$\begin{aligned} c(T) &= c'(T') && \text{since } c'_e = 0 \quad \forall e \in Z \\ &\leq c'(T^{*'}) && T' \text{ is an MCA of } G' \text{ wrt. } c'\text{-edge costs, induction hypothesis} \\ &= c'(T^*) && \text{edges of } Z \text{ have } 0 \text{ } c'\text{-cost} \end{aligned}$$

So T is an MCA of G with respect to c' -edge costs. Thus T is an MCA of G with respect to c -edge costs. \square

Matroids

4.1 Introduction

matroid

A matroid is a tuple $M = (U, \mathcal{I})$, where U is a ground set (or universe), and $\mathcal{I} \subseteq 2^U$ is a collection of subsets of U satisfying the following properties:

- (a) $\emptyset \in \mathcal{I}$.
- (b) If $A \in \mathcal{I}$, and $B \subseteq A$, then $B \in \mathcal{I}$.
- (c) (Exchange property) if $A, B \in \mathcal{I}$, with $|A| < |B|$, then $\exists e \in B - A$ such that $A \cup \{e\} \in \mathcal{I}$.

Sets in \mathcal{I} are called **independent sets**. A set not in \mathcal{I} is called a **dependent set**.

A **maximal independent set**, i.e., a set $B \in \mathcal{I}$ such that $B \cup \{e\} \notin \mathcal{I} \forall e \in U - B$, is called a **basis**.


Example: uniform matroid

Let U be any n -element set, e.g., $\{1, \dots, n\}$. Let $\mathcal{I} := \{A \subseteq U : |A| \leq k\}$, where $k \geq 0$. Then $M = (U, \mathcal{I})$ is a matroid.

- (a) $\emptyset \in \mathcal{I}$ since $|\emptyset| = 0 \leq k$
- (b) If $A \in \mathcal{I}$ and $B \subseteq A$, then $|B| \leq |A| \leq k$, so $B \in \mathcal{I}$.
- (c) Suppose $A, B \in \mathcal{I}$ with $|A| < |B| (\leq k)$. Take any $e \in B - A$, and note that $|A \cup \{e\}| = |A| + 1 \leq |B| \leq k$, so $A \cup \{e\} \in \mathcal{I}$.

Example: partition matroid

Again, let U be n -element set. Let (S_1, \dots, S_ℓ) be a partition of U , and let $r_1, \dots, r_\ell \geq 0$ be some non-negative integers. Let $\mathcal{I} := \{A \subseteq U : |A \cap S_i| \leq r_i \quad \forall i = 1, \dots, \ell\}$.



Then $M = (U, \mathcal{I})$ is a matroid.

(a), (b) hold trivially. For part (c), suppose $A, B \in \mathcal{I}$ with $|A| < |B|$. Then there exists S_i such that $|A \cap S_i| < |B \cap S_i|$ (since (S_1, \dots, S_ℓ) partition U). Consider any $e \in (B \cap S_i) - (A \cap S_i)$, and $A' = A \cup \{e\}$. Then

$$|A' \cap S_j| = \begin{cases} |A \cap S_j| \leq r_j & \text{if } j \neq i \\ |A \cap S_i| + 1 \leq |B \cap S_i| \leq r_i & \text{if } j = i \end{cases}$$

So $A' \in \mathcal{I}$.

Example:

Let U be collection of n vectors in \mathbb{R}^d . Let

$$\mathcal{I} := \{A \subseteq U : \text{the vectors in } A \text{ are linearly independent}\}$$


Then $M = (U, \mathcal{I})$ is a matroid.

(a), (b) hold trivially. (c) holds because of basic linear algebra: Suppose $A, B \in \mathcal{I}$ with $|A| < |B|$. If every $v \in B$ is a linear combinations of vectors in A , then since $|B| > |A|$, the vectors in B must be linearly dependent. Then there exists $v \in B$ such that $v \notin \text{span}(A)$ (where $\text{span}(A)$ is all vectors that are linear combinations of vectors in A). So $A \cup \{v\}$ consists of linearly independent vectors, i.e., $A \cup \{v\} \in \mathcal{I}$.

Example: graphic/cycle matroid

Let $G = (V, E)$ be an undirected graph. Let $U = E, \mathcal{I} := \{A \subseteq E : A \text{ is acyclic}\}$. Then $M = (U, \mathcal{I})$ is a matroid.

(a), (b) hold trivially. Consider (c). Suppose $A, B \in \mathcal{I}$ with $|A| < |B|$. Then $G_A = (V, A)$ has $n - |A|$ components where $n = |V|$, and $G_B = (V, B)$ has $n - |B| < n - |A|$ components.



In the picture above, we denote the components of G_B with red color. Some components of G_B can be contained within components of G_B . However, it cannot be that the vertex-set of every component of G_B is a subset of the vertex-set of some components of G_A . I.e., there exists some component of G_B that intersects at least two components of G_A as shown in the graph above.

This means there exists some $e \in B$ that connects two components of G_A . Thus $A \cup \{e\}$ is acyclic, so $A \cup \{e\} \in \mathcal{I}$.

4.2 Max-weight independent set (MWIS) problem

Here we present two matroid optimization problems.

Max-weight independent set (MWIS) problem

Given a matroid $M = (U, \mathcal{I})$, and weights $\{w_e\}_{e \in U}$ (the w_e 's could be arbitrary), find a max-weight independent set, i.e., find $A \in \mathcal{I}$ such that

$$w(A) := \sum_{e \in A} w_e = \max_{B \in \mathcal{I}} w(B)$$

Matroid intersection problem

Given two matroids $M_1 = (U, \mathcal{I}_1)$ and $M_2 = (U, \mathcal{I}_2)$ and weights $\{w_e\}_{e \in U}$, find a max-weight set that is independent in both matroids (common independent set), i.e., solve

$$\max_{A \in \mathcal{I}_1 \cap \mathcal{I}_2} w(A)$$

Matroid intersection problem is beyond the scope of this course, but you can learn it from [CO 450](#).

Note that *max-weight spanning tree* with positive edge weights is a special case of MWIS problem, where the matroid is the graphic matroid. (since bases of graphic matroid associated with a connected graph are spanning trees). MST with $\{c_e\}$ edge costs can be captured by max-weight spanning tree with positive edge weights. By defining $w_e = M - c_e$, where $M > \max_e c_e$ (so that $w_e > 0$). For any spanning tree T $w(T) = (n-1)M - c(T)$.

Now we develop the greedy algorithm for MWIS. Input is $M = (U, \mathcal{I})$, $\{w_e\}_{e \in U}$. We may assume that $w_e > 0$ for all $e \in U$. Because otherwise, we can move to the smaller matroid

$$M' = (U' := \{e \in U : w_e > 0\}, \mathcal{I}' = \{A \subseteq U' : A \in \mathcal{I}\})$$

where we can verify M' is a matroid. Solving MWIS on M' will also solve MWIS on M . Since every independent set of M' is also independent in M . If $A \in \mathcal{I}$, then $A \cap U' \in \mathcal{I}'$ and $w(A \cap U') \geq w(A)$.

Algorithm 6: Greedy Algorithm for MWIS

Input: $M = (U, \mathcal{I})$, $\{w_e\}_{e \in U}$

- 1 Sort elements in decreasing order of weight.
- 2 Initialize $A \leftarrow \emptyset$
- 3 Considering elements in sorted order, if $A \cup \{e\}$ is independent, then set $A \leftarrow A \cup \{e\}$ where e is current element being considered.
- 4 **return** A

Observe that above algorithm run on graphic matroid is equivalent to Kruskal for max-weight spanning tree.

Correctness**Claim 4.1**

If B, B' are two bases of a matroid, then $|B| = |B'|$.

Proof:

Suppose not, and $|B| < |B'|$. Then since $B, B' \in \mathcal{I}$, by the exchange property of matroids, there exists $e \in B' - B$ such that $B \cup \{e\} \in \mathcal{I}$, contradicting that B is a maximal independent set. \square

Now let A be the set returned by the greedy algorithm.

Claim 4.2

A is a basis of M .

Proof:

Suppose there exists $e \notin A$ such that $A \cup \{e\} \in \mathcal{I}$. Then consider the point when e is considered by greedy. At that point, we have some set $S \subseteq A$. But then $S \cup \{e\} \in \mathcal{I}$ since $S \cup \{e\} \subseteq A \cup \{e\}$, so algorithm should have added e , contradicting that $e \notin A$. \square

Theorem 4.3

Greedy algorithm returns a max-weight independent set.

Proof:

Let A^* be a max-weight independent set. We want to show that $w(A) = w(A^*)$, hence A is a max-weight independent set.

Observe that A^* is a basis of M . Otherwise if there exists $e \notin A^*$ such that $A^* \cup \{e\} \in \mathcal{I}$, we have $w(A^* \cup \{e\}) > w(A^*)$. So we have $|A| = |A^*|$ by Claim 4.1 and Claim 4.2.

Let $k = |A| = |A^*|$. Suppose $w(A) < w(A^*)$. Let

$$\begin{aligned} A &= \{e_1, e_2, \dots, e_k\} \\ A^* &= \{e_1^*, e_2^*, \dots, e_k^*\} \end{aligned}$$

where elements are ordered by the ordering used by greedy.

Let $A_i := \{e_1, \dots, e_i\}$ and $A_i^* := \{e_1^*, \dots, e_i^*\}$ for all $i = 1, \dots, k$. And we define $A_0 = A_0^* = \emptyset$. Consider the smallest index j such that $w(A_j) < w(A_j^*)$. Such j exists since $w(A_k) = w(A) < w(A^*) = w(A_k^*)$. We have

- $|A_j^*| = j$, $A_j^* \in \mathcal{I}$ since $A_j^* \subseteq A^*$
- $|A_{j-1}| = j-1$, $A_{j-1} \in \mathcal{I}$ since $A_{j-1} \subseteq A$
- By the exchange property, there exists $e \in A_j^* - A_{j-1}$ such that $A_{j-1} \cup \{e\} \in \mathcal{I}$.

Since $e \in A_j^*$ and e_j^* has the least weight among elements of A_j^* , we have $w_e \geq w_{e_j^*}$. Since

$$w(A_{j-1}) + w_{e_j} = w(A_j) < w(A_j^*) = w(A_{j-1}^*) + w_{e_j^*}$$

and $w(A_{j-1}^*) \leq w(A_{j-1})$. Thus $w_{e_j^*} > w(e_j)$. To summarize, we have

$$w_e \geq w_{e_j^*} > w_{e_j}$$

Consider the point when greedy considers element e . At this point, we have some set $S \subseteq A_{j-1}$. This is because $w_e > w_{e_j}$, so greedy must consider e before e_j . But then $S \cup \{e\} \in \mathcal{I}$ since $S \cup \{e\} \subseteq A_{j-1} \cup \{e\}$ and so greedy should have added e . Contradiction, since then we would have $e \in A_{j-1}$. \square

Running Time and Input Specification

How is the matroid M given as input? It would be space inefficient if we give all sets in \mathcal{I} as input. So M is specified by means of the universe U and a **matroid independence oracle**, which is a procedure that given a set $S \subseteq U$ as input, answers (correctly) if $S \in \mathcal{I}$.

Running time is the number of elementary operations + number of calls (i.e., queries) made to independence oracle. Let $m = |U|$. Then the running time is $O(m \log m) + O(m \cdot \text{oracle}) = O(m \log m)$ which is polynomial-time, i.e., efficient algorithm.

4.3 Applications of Matroid Optimization

Maximum-Weight Bipartite Matching

bipartite

A graph $G = (V, E)$ is called bipartite if V can be partitioned as $L \cup R$ such that every edge has one end in L and one end in R . We call (L, R) bipartition of V .

matching

Given a graph $G = (V, E)$, a set $M \subseteq E$ is called matching if $|M \cap \delta(v)| \leq 1$ for all $v \in V$. I.e., for every $v \in V$, there is at most one edge of M incident to it.

Max-weight bipartite matching problem

Given a bipartite graph $G = (V, E)$, edge weights $\{w_e\}_{e \in E}$, find a matching M of maximum total edge weight.

Let $L \cup R$ be bipartition of V . Define the following two matroid, having ground set E :

$$M_L = (U = E, \mathcal{I}_L = \{A \subseteq U : |A \cap \delta(v)| \leq 1, \forall v \in L\})$$

$$M_R = (U = E, \mathcal{I}_R = \{A \subseteq U : |A \cap \delta(v)| \leq 1, \forall v \in R\})$$

Because the graph is bipartite, $\{\delta(v)\}_{v \in L}$ partitions E , $\{\delta(v)\}_{v \in R}$ partitions E . Then M_L and M_R are partition matroids. Thus $A \subseteq E$ is a matching if and only if $A \in \mathcal{I}_L \cap \mathcal{I}_R$. Hence max-weight matching is equivalent to max-weight independent set in M_L and M_R . I.e., max-weight bipartite matching can be solved using matroid intersection.


Scheduling Problem

Given a set U of jobs; each $j \in U$ has **unit processing time**, deadline $d_j \geq 1$ which is integer, weights w_j (can be arbitrary). We can only process one job at any time. Find a max-weight set S of jobs, and an ordering of S such that all jobs in S complete by their deadlines.

For example,

j	1	2	3	4	5	6
d_j	3	2	1	3	2	4
w_j	2	1.8	5.6	-0.4	-1	3

Let $S = \{2, 5, 4\}$: can order S so that all jobs complete by their deadlines.



Consider $S = \{3, 2, 1, 4\}$. We cannot order S so that all jobs complete by their deadlines since there are four jobs in S with deadlines ≤ 3 .

schedulable

Say that $S \subseteq U$ is “schedulable” if there exists an order of S that completes all jobs by their deadlines.

Exercise:

If S is schedulable, then ordering jobs in S in increasing order of deadlines yields an ordering where all jobs complete by their deadlines.

S is schedulable if and only if $\forall t = 0, 1, \dots, |S|$, (number of jobs in S with deadline $\leq t$) $\leq t$

Theorem 4.4

$M = (U, \mathcal{I} = \{S \subseteq U : S \text{ is schedulable}\})$ is a matroid. Hence scheduling problem can be solved by solving MWIS for matroid M .

Proof:

Properties (a), (b) hold trivially. Consider exchange property. Let $A, B \in \mathcal{I}$ with $|A| < |B|$. We have a notation: for $S \subseteq U$, let $S_{\leq t} = \{j \in S : d_j \leq t\}$. Consider smallest $t \geq 0$ such that

$$|B_{\leq t'}| > |A_{\leq t'}| \quad \forall t' \geq t \quad (*)$$

Such a t exists, since for $t = \max_{j \in B} d_j$, $(*)$ holds and for $t = 0$, $(*)$ does not hold, and $t \geq 1$.

Claim There exists $j \in B - A$ such that $d_j = t$.

If not, then all jobs in B with deadline equal to t are also in A . Then

$$\begin{aligned} |B_{\leq t-1}| &= |B_{\leq t}| - (\# \text{ jobs in } B \text{ with deadline} = t) \\ &> |A_{\leq t}| - (\# \text{ jobs in } B \text{ with deadline} = t) && \text{due to } (*) \\ &\geq |A_{\leq t}| - (\# \text{ jobs in } A \text{ with deadline} = t) \\ &= |A_{\leq t-1}| \end{aligned}$$

Contradicts t being the smallest value such that $(*)$ holds.

Claim $A' = A \cup \{j\}$ is schedulable where $d_j = t$.

Consider any $t' = 0, 1, \dots, |A'|$. Then

$$|A'_{\leq t'}| = \begin{cases} |A_{\leq t'}| \leq t' & \text{if } t' < t \\ |A_{\leq t'}| + 1 \leq |B_{\leq t'}| \leq t' & \text{if } t' \geq t \end{cases}$$

So by exercise, A' is schedulable. So M is a matroid. □

Steiner Tree & Computational Complexity


5.1 Minimum Steiner Tree Problem

Minimum Steiner Tree Problem


Given an undirected graph $G = (V, E)$, edge costs $c_e \geq 0$, and a set $T \subseteq V$ called **terminals**, find a min-cost tree that spans (i.e., connects) all the vertices in T .

This is called Steiner tree for T and we will omit “for T ” if T is clear from the context.

Consider two examples.



Red boxes denote terminals T . Blue edges form a Steiner tree for T . Orange vertex is in Steiner tree but not in T . This is called a Steiner node. And we notice that this node is not necessary to include, but it might be useful in terms of the cost. For example,



Best Steiner tree containing only nodes from T has cost $2(|T| - 1)$. But blue edges form a Steiner tree of cost $|T|$.

Useful Transformation Let $G = (V, E)$. $(G, c, T) \rightarrow (G', c', T)$ where $G' = (V, E')$ is the complete graph on V and c'_{uv} is the shortest-path distance in G between u and v .

Note that $\forall u, v, w \in V$, $c'_{uw} \leq c'_{uv} + c'_{vw}$ (\triangle -ineq) since one u - w path in G is shortest u - v path in G + shortest v - w path in G .

Claim 5.1

- (a) Any Steiner tree F in G is also a Steiner tree in G' , and $c'(F) \leq c(F)$.
- (b) Any Steiner tree F' in G' yields a Steiner tree F'' in G such that $c(F'') \leq c'(F')$.

Proof:

For part (a), clearly F is a Steiner tree in G' . Also, note that $\forall uv \in E$, $c'_{uv} \leq c_{uv}$, and so $c'(F) \leq c(F)$.

For part (b), take F' , and “expand” each edge $uv \in F'$ to the shortest u - v path P_{uv} in G , to get an edge-set \hat{F} . (So $\hat{F} = \bigcup_{uv \in F'} P_{uv}$, where P_{uv} is the shortest u - v path in G .)

Clearly \hat{F} connects all vertices in G ; but it could contain cycles. We get F'' from \hat{F} by simply removing edges from \hat{F} so that we get an acyclic graph. So

$$c(F'') \leq c(\hat{F}) \leq \sum_{uv \in F'} c(P_{uv}) = \sum_{uv \in F'} c'_{uv} = c'(F')$$

\uparrow
 Since we removed
 edges from \hat{F} to set F''

□

(G', c') defined as above is called **metric completion** of (G, c) .

By Claim 5.1, we can always solve Steiner tree problem on the metric completion of (G, c) . The optimum values do not change, and optimum solution in the metric completion gives us back an optimal solution in the original instance and vice versa.

5.2 MST-based algorithm

Idea Find an MST in $G'[T] := (T, E'[T])$ ¹ with respect to c' -costs and map this to a Steiner tree in G using Claim 5.1 (b).

Algorithm 7: Algorithm for Minimum Steiner Tree Problems

- 1 Given instance (G, c, T) , consider the instance (G', c', T) , where (G', c') is the metric completion of (G, c) .
 - 2 Find an MST in $G'[T] := (T, E'[T])$ with respect to c' -costs and map this back to a Steiner tree in G using Claim 5.1 (b).
-

This is indeed a polytime algorithm, because all operations, including mapping back and so on, are all polynomial.

Cost of our solution $\leq \text{MST}(G', c', T)$, which c' -cost of MST in $G'[T]$. Let

$$\text{OPT} := \text{OPT}(G', c', T) = \text{OPT}(G, c, T)$$

where $\text{OPT}(G', c', T)$ is c' -cost of an optimal solution for (G', c', T) .

Theorem 5.2

(Cost of solution returned \leq) $\text{MST}(G', c', T) \leq 2 \cdot \text{OPT}$.

¹subgraph of G' induced by T . It is also complete because G is complete

Proof:

Let F^* be optimal Steiner tree for (G', c', T) . Pick some $r \in T$, root F^* at r , and do a DFS traversal of F^* starting at r .

DFS traversal yields a tour Z (i.e., a closed walk or cycle with repeated nodes) that visits all vertices of F^* , and has c' -cost $\leq 2c'(F^*) = 2 \cdot \text{OPT}$.

Will “shortcut” Z to get a cycle \hat{Z} that visits every terminal exactly once. Suppose

$$Z : u_0 = r, u_1, u_2, \dots, u_{k-1}, u_k = u_0 = r$$

$$\hat{Z} : u_0 = r, u_{i_1}, u_{i_2}, \dots, u_{i_{\ell-1}}, u_{i_\ell} = u_0 = r$$

where for every j , u_{i_j} is the first terminal after $u_{i_{j-1}}$ in the tour Z that has not been visited before.


By \triangle -ineq, or the definition of c' , we have $c'_{u_{i_j}u_{i_{j+1}}} \leq c'(Z_{u_{i_j}u_{i_{j+1}}})$ where $Z_{u_{i_j}u_{i_{j+1}}}$ is the portion of Z between first occurrences of u_{i_j} and $u_{i_{j+1}}$.

Since $Z_{u_{i_j}u_{i_{j+1}}}$ are disjoint for different j we have $c'(\hat{Z}) \leq c'(Z)$. Removing an edge from \hat{Z} gives a spanning tree in $G'[T]$, of cost $\leq c'(\hat{Z}) \leq c'(Z) = 2 \cdot \text{OPT}$.

Thus $\text{MST}(G', c', T) \leq 2 \cdot \text{OPT}$. □

Example:

Here is an example of how we get \hat{Z} .




Terminal nodes are denoted by blue. So DFS traversal gives us Z

$$r, 1, 2, 1, r, 3, 4, 5, 4, 6, 4, 3, 7, 3, r, 8, 9, 8, 10, 8, r$$

We mark first occurrences of terminals in Z . And these marked vertices form \hat{Z} .

Were we dumb in analyzing this algorithm? Actually no. Consider an example.



where $\epsilon > 0$ is small. Our algorithm returns all edges of cost 2ϵ , thus the total cost is $(2 - \epsilon)(|T| - 1)$.

But the optimal solution is blue edges, of cost $|T|$. We see the factor is approximately 2. Can we design a better algorithm? For this problem, there is no known efficient algorithm that is always guaranteed to return an optimal solution. Moreover, researchers don't believe such an algorithm exists.

α -approximation algorithm

An α -approximation algorithm, where $\alpha \geq 1$ is the approximation factor, for a minimization problem, is an efficient (i.e., polytime) algorithm that on every instance returns a solution of cost $\leq \alpha \cdot \text{OPT}$.

So above algorithm for Steiner tree is a 2-approximation algorithm for the Steiner tree problem.

5.3 Computational Complexity

P is class/set of all problems that can be solved by polytime algorithms. For example, Shortest Paths, MST, MCA $\in P$. Other YES/NO problems:

- **IsComposite**: Given integer $n \geq 2$, is n a composite number; i.e, do exist integers $x, y \geq 2$ such that $n = xy$?
- **Factoring**: Given integer $n \geq 2$, decide if n is composite and if so, find integers $x, y \geq 2$ such that $n = xy$.
- **IsPrime**: Given integer $n \geq 2$, is n a prime number?

Algorithm 8: Factoring algorithm: Sieve of Eratosthenes

- 1 Consider all integers x such that $2 \leq x \leq \sqrt{n}$, and see if x divides n .
 - 2 If so, then n is composite and $x, \frac{n}{x}$ is a factorization of n .
-

Is the above a polytime algorithm? No, since number of iterations $\approx \sqrt{n}$, which is not polynomial in $O(\log n)$ which is the input size.

Open question: Is **Factoring** $\in P$? We already know that **IsPrime** $\in P$, which is only settled in around 2003. This implies **IsComposite** $\in P$.

Decision problem is problem with a YES/NO answer. We can take any optimization problem and consider a decision-version of the problem:

- Is there a solution of cost $\leq k$ (for minimization problems)
- Is there a solution of value $\geq k$ (for maximization problems)

Here is decision version of MST (**DecMST**): given $(G, \{c_e\})$, and a number k , is there a spanning tree of cost $\leq k$? If **MST** $\in P$, then **DecMST** $\in P$.

5.4 Polynomial-Time Reductions

polytime reduction

Given problems A, B , we say A **reduces in polynomial time to B** , or A is **polytime reducible to B** , denote $A \leq_p B$, if we can solve problem A using a polynomial number of elementary operations + a polynomial number of calls to an algorithm for problem B .

$MST \leq_p MCA$. Given input $(G = (V, E))$ undirected graph, $\{c_e\}_{e \in E}$ to MST problem, create $(\overset{\leftrightarrow}{G}, \overset{\leftrightarrow}{c})$ the bidirected version of G , by creating edges $(u, v), (v, u)$ in $\overset{\leftrightarrow}{G}$ for every edge uv of G , and give these two edges cost c_{uv} . Solve MCA on $(\overset{\leftrightarrow}{G}, \overset{\leftrightarrow}{c}, r)$ where r is any node of V , and map back MCA to undirected edge of G to get an MST.

$MST \leq_p$ Min Steiner Tree Problem. This is true because MST is a special case of Min Steiner Problem, when edge costs are ≥ 0 . For MST, can always ensure (i.e., move to) nonnegative edge costs by adding a suitable large constant to all edge costs.

Suppose $A \leq_p B$. Then

1. If $B \in P$, then $A \in P$.
2. Equivalently, if $A \notin P$, then $B \notin P$.

Consider

- SORTING: Given n numbers a_1, \dots, a_n , sort them in an increasing order.
- MIN: Given n numbers a_1, \dots, a_n , find the minimum of these n numbers.

$SORTING \leq_p MIN$: Can sort using n calls to an algorithm for MIN. Also $MIN \leq_p SORTING$.

Consider ²

- LPFeas: Given $A \in \mathbb{Q}^{m \times n}$, is the system $Ax \leq b, x \geq 0$ feasible?
- LP0pt: Given $A \in \mathbb{Q}^{m \times n}, b \in \mathbb{Q}^m, c \in \mathbb{Q}^n$, does the LP

$$\begin{aligned} \max \quad & c^T x \\ \text{s.t.} \quad & Ax \leq b \\ & x \geq 0 \end{aligned}$$

have an optimal solution?

$LPFeas \leq_p LP0pt$. We can set $c = 0$, and call algorithm for LP0pt. Using duality we can show $LP0pt \leq_p LPFeas$.

Decision problem LPImprov: Given $A \in \mathbb{Q}^{m \times n}, b \in \mathbb{Q}^m, c \in \mathbb{Q}^n$, and $z \in \mathbb{Q}$, determine if there is a feasible solution to the LP:

$$\begin{aligned} \max \quad & c^T x \\ \text{s.t.} \quad & Ax \leq b \\ & x \geq 0 \end{aligned}$$

with objective value $> z$? Using duality, can show that $LPImprov \leq_p LPFeas$.

Note that simplex method for solving LPs is NOT a polytime algorithm.

Consider

- MST: Given connected graph $G = (V, E)$, edge costs $\{c_e\}$ that are positive integers, find cost of MST.
- Decision version of MST, DecMST: Given connected $G = (V, E)$, positive integer edge costs $\{c_e\}$, and an integer $k > 0$, is there a spanning tree of cost $\leq k$?

We observe that $DecMST \leq_p MST, MST \leq_p DecMST$.

Given MST instance (G, c) , we can deduce that every spanning tree has cost $\geq LB := 1$, and MST has cost $\leq UB := (n - 1)c_{\max}$, where $c_{\max} := \max_e c_e$.

²Here we choose to work with rational numbers because it's hard to specify the input size for irrational numbers

Algorithm Start at $k = UB$. Solve DecMST with this k ; if answer is YES, set $k \leftarrow k - 1$ and REPEAT.

This is not a polytime reduction, since we could be calling algorithm for DecMST $\approx UB$ times, but $UB \approx (n - 1)c_{\max}$ is not polynomially bounded in input size.

We can get polynomial reduction by doing binary search. Binary search calls algorithm DecMST $O(\log_2 UB)$ which is polynomially bounded in input size.

It's not hard to see that \leq_p is transitive.

Consider two problems:

- MST: find cost of MST
- A-MST: Find a min-cost spanning tree.

We can see that $MST \leq_p A\text{-}MST$. $A\text{-}MST \leq_p MST$: Let $k = MST(G, c)$. Let e be some edge of G . Then $k' = MST(G - e, \{c_f\}_{f \in G - e})$. If $k' > k$, then e belongs to every MST. If $k' \leq k$, then $G \leftarrow G - e$, $c \leftarrow \{c_f\}_{f \neq e}$.

P & NP

Consider decision version of Steiner tree:

DecSteiner: Given an instance (G, c, T) and a number k , is there a Steiner tree of cost $\leq k$?

Observation If DecSteiner instance is a YES instance, then there is a simple certificate (A Steiner tree of cost $\leq k$) to convince one that answer is indeed YES.

verifier/certifier

A verifier/certifier V for a decision problem Π is an algorithm that takes two inputs, x : an instance of Π , and a “certificate” y , and outputs YES or NO. It satisfies

- If x is a YES instance, then there exists y such that $V(x, y) = \text{YES}$.
- If x is a NO instance, then for all y , $V(x, y) = \text{NO}$.

6.1 NP

NP a decision problem Π is in the class **NP** if there exists a polynomial p and a polytime verifier V for Π such that

- For every YES instance x of Π , there exists a certificate y with $\text{size}(y) \leq p(\text{size}(x))$ such that $V(x, y) = \text{YES}$.
- (For every NO instance x , $V(x, y) = \text{NO}$ for all y)

Informally, a decision problem is in **NP** if its YES instances have “short”, efficiently verifiable certificates.

Example:

DecSteiner \in **NP**. It has certificate y : Steiner tree of cost $\leq k$. Verifier V : check y is indeed a Steiner tree of cost $\leq k$.

Example:

Recall **IsComposite**. It is in **NP**. Certificate y : a factor of n where $2 \leq y < n$. Verifier: check y divides n (and $2 \leq y < n$).

Example:

$\text{IsPrime} \in NP$? Yes, certificate for YES instance uses results in number theory, Fermat's little theorem. We know that $\text{IsPrime} \in P$.

Claim 6.1

If Π is a decision problem in P , then Π is also NP . In other words, $P \subseteq NP$.

Proof:

Let A : polytime algorithm for Π . Then, we can construct a polytime verifier V for Π as follows: on input (x, y) , V returns $A(x)$. \square

Exercise:

Prove that $P \neq NP$.

Example: 3-SAT

Given a boolean formula F of the form

$$F = C_1 \wedge C_2 \wedge \cdots \wedge C_m,$$

where each C_i is a clause of the form $y_{i_1} \vee y_{i_2} \vee y_{i_3}$ where each y_i (called literal) is x_i or \bar{x}_i .

Is there a True/False assignment of the variable under which F evaluates to T? Such an assignment is called a satisfying assignment.

Not hard to see $3\text{-SAT} \in NP$: certificate y : satisfying assignment. Verifier: check that y is indeed a satisfying.

Example: Set-Cover

Input: a universe U , and a collection $\mathcal{S} = \{S_1, \dots, S_m\}$ of subsets of U , and an integer $k \geq 0$.

Are there k subsets $S_{i_1}, \dots, S_{i_k} \in \mathcal{S}$ such that $(S_{i_1} \cup \cdots \cup S_{i_k}) = U$?

$\text{Set-Cover} \in NP$: a certificate for YES instance is simply k sets of \mathcal{S} whose union is U .

6.2 NP-hard and NP-complete

NP-hard, NP-complete

A problem B is called:

- NP-hard: if $X \leq_p B \forall x \in NP$
- NP-complete: if $B \in NP$ and B is NP-hard (i.e., $x \leq_p B \forall x \in NP$)

Intuitively, NPC are hardest problems in class NP .

Suppose Y is a NPC problem. Then

- if $Y \in P$, then $P = NP$. Since $X \leq_p Y$ for all $X \in NP$ and $Y \in P$, we get that $x \in P \forall x \in NP$.
- If $Y \notin P$, then $P \neq NP$.

How do we show that a problem Y is NP-hard or NPC?

- Pick a suitable NP-hard problem B .
- Show $B \leq_p Y$. Then Y is NP-hard by transitivity of \leq_p .

- If we want to show that Y is NP-complete, then also SHOW $Y \in NP$, which is usually easy.

Note:

If Y' is the decision version of optimization problem Y and Y' is NP-hard, then Y is also NP-hard.
Since $Y' \leq_p Y$.

NP-Hard and NP-Complete

Theorem 7.1: Cook-Levin Theorem

3-SAT is NP-complete.

Will show $\Pi_{\text{set cover}}$ is NPC by showing that $3\text{-SAT} \leq_p \text{Set Cover}$. (Already know $\text{set cover} \in NP$)

Will show that $3\text{-SAT} \leq_p$ special case of set cover, where we have an undirected graph $G = (V, E)$, and that defines

- $U = E$
- every vertex $v \in V$ creates a set in \mathcal{S} given by $\delta(v)$.

So a collection of sets $\mathcal{S}' = \{\delta(v) : v \in T\}$ is a set cover. Here T is called vertex cover.


Theorem 7.2

$3\text{-SAT} \leq_p \text{Vertex Cover}$.

Proof:

Let $F = C_1 \wedge C_2 \wedge \dots \wedge C_m$ where each $C_i = y_{i1} \vee y_{i2} \vee y_{i3}$ be an instance of 3-SAT. Will create a VC-instance such that F is satisfiable \Leftrightarrow VC instance is a YES instance.

- For every clause $C_i = y_{i1} \vee y_{i2} \vee y_{i3}$, create the triangle,




- For every var x_j if $x_j \in C_i$, and $\bar{x}_j \in C_k$, then we create edge $(x_j, i) - (\bar{x}_j, k)$.

So graph $G = (V, E)$ has

$$V = \{(y_j, k) : y_j \in C_k\} \quad E = \{(x_j, i)(x_{j'}, k) : i = k, \text{ or } x_j = \bar{x}_{j'}\}$$

For example, let $F = (x_1 \vee \bar{x}_2 \vee x_3) \wedge (\bar{x}_1 \vee x_4 \vee x_2) \wedge (x_3 \vee x_4 \vee \bar{x}_1)$.



We want to show that F is satisfiable \Leftrightarrow the graph created has a VC of size (number of nodes in VC solution) $2m$ where m is number of clauses in F .

Suppose F satisfiable. Then for every C_i , at least one literal in C_i is set to True. Pick exactly one literal, say $y_i^* \in C_i$, that is set to True and consider set $S = \{(y_i, i) : y_i \neq y_i^*\}$. So $|S| = 2m$. We claim that S is a vertex cover.

Clearly all \triangle -edges are covered since we pick 2 nodes from each \triangle . Consider an edge $(x_j, i)(\bar{x}_j, k)$. It cannot be that both x_j, \bar{x}_j are true, so at least one of these has not been picked for the corresponding clauses (as y_i^*), and so is in S .

Conversely, **suppose the graph has been created has a VC S with $|S| \leq 2m$.** Note that $|S| = 2m$, since S must contain ≥ 2 nodes from each triangle, therefore S contains exactly 2 nodes from C_i 's clause- \triangle .

Suppose for C_i , S contains nodes $(y_{i_1}, i), (y_{i_2}, i)$. Then we set $y_{i_3} = \text{true}$. We claim that this gives a satisfying (possibly partial) assignment. Clearly, every clause is set to be true. Consider some var x_j and suppose $(x_j, i)(\bar{x}_j, k) \in V$. Then at least one of these nodes is in S , so we would not set both x_j, \bar{x}_j to true, which is then a valid assignment as this holds for every variable. \square


Theorem 7.3

Set Cover \leq_p Steiner Tree.

These are decision versions of the problems. Hence, decision version of Steiner tree is NP-complete.

Proof:

Let $(U, S \subseteq 2^U, k)$ be a set cover instance. We represent SC instance (U, S) as follows:



We create Steiner tree instance by appending bipartite graph representing (U, S) with a “root”

node r , and edges $rs \forall s \in \mathcal{S}$. Terminal set $T = \{r\} \cup U$. Set edge cost

$$c_f = \begin{cases} 1 & f \in \delta(r) \\ k+1 & f \text{ is incident to an element of } U \end{cases}$$

Then we want to show \exists set cover of size $\leq k \Leftrightarrow \exists$ Steiner tree of cost $\leq \underbrace{n(k+1)}_M + k$ where $n = |U|$.

Suppose $\mathcal{S}' \subseteq \mathcal{S}$ is a collection of $\leq k$ sets whose union is U , the tree with edges $\{rs_i : S_i \in \mathcal{S}'\}$ and $\{eS_j : S_j \text{ in some set in } \mathcal{S}' \text{ containing } e\}$ is a Steiner tree of cost $\leq n(k+1) + k$.

Suppose F is a Steiner tree of cost $\leq M$. Then every node $e \in U$ must be a leaf node in F , otherwise $c(F) \geq (k+1)(\# \text{ edges in } F \text{ incident to nodes in } U) \geq (n+1)(k+1) > M$. Then there are n edges joining edges $e \in U$ to set nodes. Then there are $\leq k$ edges from $\delta(r)$ are in F .

Consider $\mathcal{S}' = \{S_i : rs_i \in F\}$. Then $|\mathcal{S}'| \leq k$ and every $e \in U$ is connected to r by a 2-hop path $e-S_i-r$, and so $S_i \in \mathcal{S}'$. Therefore \mathcal{S}' is a set cover. \square


Approximation Algorithm & Primal Dual Algorithm

8.1 Approximation Algorithm Design

Recall the definition for α -approximation algorithm where $\alpha \geq 1$.

Strategy for approximation algorithm design

1. Come up with a (move tractable) lower bound LB on OPT ($LB \leq OPT$)
2. Come up with a polytime algorithm and show that it returns a solution of cost $\leq \alpha \cdot LB \leq \alpha \cdot OPT$.



How to come up with lower bounds? We can formulate discrete optimization problem as an integer program, and relax integrality constraints to get linear program. Then $OPT_{LP} \leq OPT_{IP}$ and now OPT_{LP} is LB.

Example: Steiner tree

Instance $(G = (V, E), T \subseteq V, \{c_e \geq 0\}_{e \in E})$

Variables: $x_e \in \{0, 1\} \forall e \in E$, and $x_e = 1$ indicates e is in opt tree; $x_e = 0$ indicates e is NOT in our tree. Here x_e is binary indicator variables.

Then the IP:

$$\begin{aligned} \min \quad & \sum_{e \in E} c_e x_e \\ \text{s.t.} \quad & x_e \in \{0, 1\} \quad \forall e \in E \\ & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \subseteq V : S \cap T \neq \emptyset, T - S \neq \emptyset \end{aligned} \quad (\text{IP})$$

The second condition means for every cut $(S, V - S)$ where S and S^C contain some terminals, then the boundary of the cut must include at least one edge from the tree, otherwise tree might not connect up all terminal nodes.

We then can relax first integrality constraint to $0 \leq x_e \leq 1$, or even further $x_e \geq 0$ because all

feasible $x_e > 1$ can be “reduced to” $x_e = 1$, which still satisfies the constraints. Then LP is

$$\begin{aligned} \min \quad & \sum_{e \in E} c_e x_e \\ \text{s.t.} \quad & x_e \geq 0 \quad \forall e \in E \\ & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \subseteq V : S \cap T \neq \emptyset, T - S \neq \emptyset \end{aligned} \quad (\text{LP})$$

Then $OPT_{LP} \leq OPT_{IP}$.

LP relaxation will be exploited in two ways for designing algorithms

1. LP-rounding algorithms: Solve LP-relaxation to get a (likely) fractional solution x^* . Devise a polytime algorithm to transform x^* to an integer solution \hat{x} and show that

$$\text{cost}(\hat{x}) \leq \alpha \cdot \text{cost}(x^*) = \alpha \cdot OPT_{LP} \leq \alpha \cdot OPT$$

And this way of transforming is known as LP-rounding algorithm.

Theorem 8.1

LPs can be solved in polytime. (BUT simplex is NOT polytime)

Even there are exponentially many constraints, we can still solve LP in polytime via ellipsoid. Interested readers can consult CO 255 or CO 471.

2. Primal-Dual Algorithms. Recall dual of (LP) is a maximization LP (D), and every dual feasible solution y has value at most OPT_{LP} . So if we design an algorithm that constructs an integer feasible solution \hat{x} and a dual solution y , and show that $\text{cost}(\hat{x}) \leq \alpha \cdot (\text{value}(y))$. Then we get an α -approximation algorithm since $\text{value}(y) \leq OPT_{LP} \leq OPT$.

Recap of LP duality It's bad to memorize the table as said by the instructor. It's better to understand the logic of taking dual.

Back to Steiner tree...

$$\begin{aligned} \min \quad & \sum_{e \in E} c_e x_e \\ \text{s.t.} \quad & x_e \geq 0 \\ & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \in \mathcal{S} \end{aligned} \quad (\text{P})$$

where $\mathcal{S} = \{S \subseteq V : S \cap T \neq \emptyset, T - S \neq \emptyset\}$.

Now consider the dual constraint:

$$\sum_{e \in E} x_e \left(\sum_{\substack{S \in \mathcal{S}: \\ e \in \delta(S)}} y_S \right) = \sum_{S \in \mathcal{S}} y_S \left(\sum_{e \in \delta(S)} x_e \right) \geq \sum_{S \in \mathcal{S}} y_S \quad (**)$$

To ensure $\sum_{e \in E} c_e x_e \geq \text{LHS of } (**)$, we need that $c_e \geq \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S$. So the dual of (P) is

$$\begin{aligned} \max \quad & \sum_{S \in \mathcal{S}} y_S \\ \text{s.t.} \quad & \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S \leq c_e \quad \forall e \in E \\ & y \geq 0 \end{aligned} \quad (\text{D})$$

Weak Duality

The value of any feasible solution y to (D) provides a lower bound on $OPT(P)$ (when (P) is minimization LP).

So if x is feasible solution to (P), y feasible to (D), then $\sum_{e \in E} c_e x_e \geq \sum_{S \in \mathcal{S}} y_S$.

Strong Duality

If (P) and (D) are both feasible, then they both have optimal solutions, and $\text{OPT(P)} = \text{OPT(D)}$.

So $\sum_{e \in E} c_e x_e^* = \sum_{S \in \mathcal{S}} y_S^*$.

x^* and y^* optimal \Leftrightarrow

- x^*, y^* feasible, AND
- x^*, y^* satisfy the following Complementary Slackness conditions:

$$1. x_e^* > 0 \Rightarrow \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S^* = c_e$$

$$2. y_S^* > 0 \Rightarrow \sum_{e \in \delta(S)} x_e^* = 1$$

8.2 Network Connectivity Problems**Modeling Framework**

We will model a network connectivity problem by a cut-requirement function $f : 2^V \rightarrow \{0, 1\}$ where $f(S) = 1$ for $S \subseteq V$ indicates that a feasible solution and must include an edge from $\delta(S)$. This gives rise to the following **f -connectivity problem**: Find a min-cost set of edges F s.t. $F \cap \delta(S) \neq \emptyset \forall S \subseteq V$ s.t. $f(S) = 1$.

Example: Steiner tree

Steiner tree is modeled by the cut-requirement function f^{ST} where

$$f^{ST}(S) = \begin{cases} 1 & \text{if } S \cap T \neq \emptyset, T - S \neq \emptyset \\ 0 & \text{otherwise} \end{cases}$$

Then Steiner tree is equivalent to f^{ST} -connectivity problem. I.e., F a Steiner tree $\Leftrightarrow F$ is an acyclic solution to f^{ST} -connectivity.

Note that $f(S) = 0$ does not say anything about $F \cap \delta(S)$.

8.3 LP-relaxation for f -connectivity and its dual

Define $\mathcal{S} = \{S \subseteq V : f(S) = 1\}$.

$$\begin{aligned} \min \quad & \sum_{e \in E} c_e x_e \\ \text{s.t.} \quad & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \in \mathcal{S} \\ & x \geq 0 \end{aligned} \tag{P}$$

$$\begin{aligned} \max \quad & \sum_{S \in \mathcal{S}} y_S \\ \text{s.t.} \quad & \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S \leq c_e \quad \forall e \in E \\ & y \geq 0 \end{aligned} \tag{D}$$

We will consider f -connectivity for a structured class of cut-requirement functions.

$\{0,1\}$ -proper cut-requirement function

A $\{0,1\}$ -proper cut-requirement function is a function $f : 2^V \rightarrow \{0,1\}$ such that

- (i) $f(V) = 0$;
- (ii) $f(S) = f(V - S) \forall S \subseteq V$; (symmetry)
- (iii) for any $A, B \subseteq V$ s.t. $A, B \neq \emptyset, A \cap B = \emptyset, f(A \cup B) = 1 \Rightarrow f(A) = 1$ or ^a $f(B) = 1$.


^anon-exclusive, could have $f(A) = f(B) = 1$

We observe that f^{ST} is a $\{0,1\}$ -proper function.

Example: Generalized Steiner tree/Steiner forest

Given undirected graph $G = (V, E)$, edge costs $\{c_e \geq 0\}$, k -disjoint terminal sets $T_1, \dots, T_k \subseteq V$, find a min-cost set of edges F s.t. each $T_i, i = 1, \dots, k$, belongs to a connected components of (V, F) in T_i . I.e., all nodes are connected in F .

For example,



where four colors represent four T_i 's, and the purple forest is a feasible solution.

This can be modelled by the proper function.

$$f^{GST}(S) = \begin{cases} 1 & \text{if } \exists i \text{ such that } S \cap T_i \neq \emptyset, T_i - S \neq \emptyset \\ 0 & \text{otherwise} \end{cases}$$

It's not hard to show that Generalized Steiner tree is equivalent to f^{GST} -connectivity: write one T_i , Generalized Steiner tree \rightarrow Steiner tree, $f^{GST} \rightarrow f^{ST}$.

We claim that f^{GST} is a proper function.

Proof:

Clearly $f^{GST}(V) = 0$ and symmetry holds. For (iii), suppose $A, B \neq \emptyset, A \cap B = \emptyset$, s.t. $f^{GST}(A \cup B) = 1$. So there exists i such that $(A \cup B) \cap T_i \neq \emptyset, T_i - (A \cup B) \neq \emptyset$. I.e.,

- (a) $A \cap T_i \neq \emptyset$ (i.e., $T_i - A \neq \emptyset$) OR
- (b) $B \cap T_i \neq \emptyset$ (i.e., $T_i - B \neq \emptyset$).

So $f^{GST}(A) = 1$ (a) OR $f^{GST}(B) = 1$ (b). □

Goal Design a 2-approximation algorithm for f -connectivity with $\{0,1\}$ -proper function.

Recall (P) and (D) defined previously. **Primal-Dual Algorithm** will simultaneously construct an integer feasible solution to (P), and a dual feasible solution y .

It's useful to introduce a notion of "time" t . Initially $t = 0$ and t increases at rate 1. (t is not necessarily an integer) We will maintain a set F of edges, and a dual feasible solution y . Initially $F \leftarrow \emptyset, y_S \leftarrow 0 \forall S \in \mathcal{S}$, which is feasible. At any state, given *infeasible* solution F , there will be some **violated sets**,

i.e., set S with $f(S) = 1$ BUT $F \cap \delta(S) = \emptyset$. We will increase y_S for some violated sets.

Which violated sets do we pick?

Design rule 1 Will pick the (inclusion-wise) minimal violated sets (MVSs). Let

$$\mathcal{V} = \{S \in \mathcal{S} : S \text{ is violated, and } S \text{ is minimal among violated sets}\}$$

“Minimal among violated sets”: $\forall T \subsetneq S, T$ is not violated.

Now given \mathcal{V} , we want increase $y_S \forall S \in \mathcal{V}$.

Lemma 8.2

Given set F of edges, the MVSs are $\{S \subseteq V : S \text{ is a component of } (V, F) \text{ with } f(S) = 1\}$.

Proof:

Clearly if $F \cap \delta(S) \neq \emptyset$, then S is not violated. So any violated set S must be a union $C_1 \cup \dots \cup C_k$ of components of (V, F) . Since $f(S) = 1$, by property (iii) of proper functions, we must have $f(C_i) = 1$ for some $i = 1, \dots, k$. Thus C_i is a MVS. \square

Primal Dual Algorithm cont'd

9.1 Primal dual algorithm for f -connectivity

Recall $\mathcal{S} = \{S \subseteq V : f(S) = 1\}$.

$$\begin{array}{ll} \min & \sum_{e \in E} c_e x_e \\ \text{s.t.} & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \in \mathcal{S} \\ & x \geq 0 \end{array} \quad (\text{P})$$

$$\begin{array}{ll} \max & \sum_{S \in \mathcal{S}} y_S \\ \text{s.t.} & \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S \leq c_e \quad \forall e \in E \\ & y \geq 0 \end{array} \quad (\text{D})$$

Corollary 9.1

$F \subseteq E$ is feasible (for f -connected with a $\{0,1\}$ -proper function) $\Leftrightarrow f(S) = 0$ for every component S of (V, F) .

Design Rule 2 Increase $y_S \forall S \subseteq \mathcal{V}$ uniformly at rate 1, (but y_S 's need not be integral) until constraint $\sum_{S \in \mathcal{S}: e \in \delta(S)} y_S \leq c_e$ goes tight for some edge $e \in \delta(S)$ for some MVS S . When this happens, we add e to F .

We STOP, when F is feasible, i.e., there are no violated sets (and hence no MVSs), where the feasibility of F can be detected using Corollary 9.1.

Final Design Rule (REVERSE DELETE) After we have stopped adding edges, consider edges in F in

reverse insertion order and delete an edge e if $F - \{e\}$ is feasible.

Algorithm 9: Primal Dual Algorithm

- 1 Initialize $F \leftarrow \emptyset$, $y_S \leftarrow 0 \forall S \in \mathcal{S}$, and time $t \leftarrow 0$ (t increase at rate 1 throughout).
Let \mathcal{V} : MVSs give F , i.e., $\{S \subseteq V : S \text{ is a component of } (V, F) \text{ with } f(S) = 1\}$. // Lemma 8.2
 - 2 **while** $\mathcal{V} \neq \emptyset$ **do**
 - (a) Increase y_S uniformly at rate 1 $\forall S \in \mathcal{V}$ until some edge $e \in \delta(S)$ goes tight for some $S \in \mathcal{V}$.
 - (b) Update $F \leftarrow F \cup \{e\}$, update \mathcal{V} .
 - 3 (REVERSE DELETE) Let $F = \{e_k, e_{k-1}, \dots, e_1\}$ where e_k is inserted last, e_{k-1} before e_k and so on.
For $i = k..1$, if $F - \{e_i\}$ is feasible, $F \leftarrow F - \{e_i\}$.
 - 4 **return** F
-

Example:


Consider Steiner tree with $T = \{s, t\}$, s - t shortest path. At each point of time, we have exactly 2 MVSs, consisting of components of (V, F) are containing s , other containing t . Then Primal Dual algorithm is equivalent to bidirectional version of Dijkstra, where we grow explored set simultaneously from both s, t .

Example:


Consider Steiner tree with $T = V$, i.e., MST. Then Primal Dual algorithm is exactly Kruskal's algorithm.

First edge to tight is least-cost edge. In general, given F , next edge to go tight in the least cost edge joining 2 components of (V, F) . This is because any edge e that gets added to F by Primal Dual algorithm, is on the boundary of exactly 2 sets in \mathcal{V} at all times until it is added. Thus e gets added at time $t_e = c_e/2$.


Now consider the example in Generalized Steiner Tree. Let $T_i = \{s_i, t_i\}$ for $i = 1, 2, 3, 4$.




Initially, all singleton nodes are violated sets as $f(\{v\}) = 1$ and it requires an edge in $\delta(\{v\})$ for all $v \in \cup T_i$.



$t = 0$, $y_S \leftarrow 0 \forall S \in \mathcal{S}$, $F \leftarrow \emptyset$.




$t = 0.5$, $y_{\{v\}} = 0.5 \forall v \in \cup T_i$. $F \leftarrow \{s_4 t_4, s_3 t_1, s_2 t_2\}$. $\mathcal{V} \leftarrow \left\{ \{s_1\}, \{s_3, t_1\}, \{t_3\} \right\}$. The red values are y_S .
 $F = \{s_2 t_2, s_3 t_1, s_4 t_4\}$.




$t = 1$, s_3t_3 goes tight. $y_{\{s_3, t_1\}} = 0.5$, $y_{\{s_1\}} = y_{\{t_3\}} = 1$, $y_{\{v\}} = 0.5 \forall v \in \cup T_i - \{s_1, t_3\}$.


Here $\mathcal{V} = \left\{ \underset{1}{\{s_1\}}, \underset{0}{\{s_3t_1t_3\}} \right\}$. $F = \{s_2t_2, s_3t_1, s_4t_4, s_3t_3\}$.



$t = 2.5$, t_1s_4 goes tight. Here $\mathcal{V} = \left\{ \underset{2.5}{\{s_1\}}, \underset{0}{\{s_3s_4t_1t_3t_4\}} \right\}$. $F = \{s_2t_2, s_3t_1, s_4t_4, s_3t_3, t_1s_4\}$.



$t = 3.5$, s_1s_2 goes tight. $\mathcal{V} = \left\{ \underset{0}{\{s_1s_2t_2\}}, \underset{1}{\{s_3s_4t_1t_3t_4\}} \right\}$. $F = \{s_2t_2, s_3t_1, s_4t_4, s_3t_3, t_1s_4, s_1s_2\}$.



$t = 5.5$, s_2s_3 goes tight. Final dual solution is

$y_{s_1} = 3.5$	$y_{t_1} = 0.5$	$y_{s_2t_2} = 0$	$y_{s_1s_2t_2} = 2$
$y_{s_2} = 0.5$	$y_{t_2} = 0.5$	$y_{s_3t_1} = 0.5$	$y_{s_3s_4t_1t_3t_4} = 3$
$y_{s_3} = 0.5$	$y_{t_3} = 1$	$y_{s_4t_4} = 0$	
$y_{s_4} = 0.5$	$y_{t_4} = 0.5$	$y_{s_3t_1t_3} = 1.5$	

And there's no violated sets!

Then the final edge set F is



where $t_1 s_4$ gets deleted.

Here $\{s_2, t_2\}, \{s_4, t_4\}$ are components that are not violated, i.e., $f(C) = 0$.

9.2 Analysis

Now back to the algorithm. We want to analyze **running time**. We only need to keep track of non-zero y_S values. We can easily identify MVSs, $\leq n$ MVSs any point.

F is always acyclic, since we only add edges on boundaries of some components.¹ Thus $|F| \leq n - 1$. Therefore, at most $n - 1$ iterations of WHILE loop in step 2. Moreover, as there are at most n MVSs, then there are at most n^2 sets S with $y_S > 0$.

We can compute (in polytime) next edge to go tight (and also the time at which this happens) by keeping track of “residual cost” $c_e - \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S$ for each $e \in E$. This implies WHILE loop runs in polytime, and REVERSE DELETE takes polytime.

Now let's examine the **approximation guarantee**. Recall the CS conditions for (P), (D):

1. $x_e > 0 \Rightarrow \sum_{S \in \mathcal{S}: e \in \delta(S)} y_S = c_e \forall e \in E$. This holds by design $\forall e \in F$.
2. $y_S > 0 \Rightarrow \sum_{e \in \delta(S)} x_e = 1 \forall S \in \mathcal{S}$. This can be violated.

Exercise:

If (1) holds, and the following relaxation of (2) holds: (2') $y_S > 0 \Rightarrow \sum_{e \in \delta(S)} x_e \leq 2$ (where x is $\{0, 1\}$ -solution corresponding to F : $x_e = 1 \forall e \in F$, $x_e = 0$ otherwise).

Then $\sum_e c_e x_e \leq 2 \sum_{S \in \mathcal{S}} y_S$, which is then a 2-approximation.

Notation For an edge set $Z \subseteq E$, $\delta_Z(S) = \delta(S) \cap Z$. From now on, F is the **final-set** of edges returned.


Lemma 9.2

At any point in the algorithm, $\sum_{S \in \mathcal{V}} |\delta_F(S)| \leq 2|\mathcal{V}|$.

This is like an “averaged” version of (2').


Proof:

Recall F is final edge set. Consider (start of) some stage i . Let $F' \supseteq F$ be the edge-set BEFORE REVERSE DELETE. Let $F'_i \subseteq F'$ be edges added to F' by START of stage i . Let $\mathcal{V} = \mathcal{V}_i$: collection of MVSs at start of stage i .



¹A tiny note: we might encounter the situation that at some point, two edges connecting two components go tight at the same time. We are processing them in linear order. After we add one edge, then the other edge is no longer on the boundary, so the other edge is not added.

Let H be graph obtained from (V, F) by contracting the node-sets corresponding to components of (V, F'_i) where we eliminate self loops.



For each node u in H , let $S_u \subseteq V$ be nodes that were contracted to form u . So $\delta_F(S_v) = \delta_H(v)$. We have to show

$$\sum_{v \in H: S_v \in \mathcal{V}} |\delta_H(v)| \leq 2|\{v \in H : S_v \in \mathcal{V}\}|$$

We claim that H is acyclic.

Consider H' be graph obtained from (V, F') by contracting components of (V, F'_i) . Since $F', F'_i \subseteq F'$ are acyclic, H' is acyclic. But H is a subgraph of H' , so H is acyclic.


Consider a component C of H . We will show

$$\sum_{v \in C: S_v \in \mathcal{V}} |\delta_H(v)| \leq 2|\{v \in C : S_v \in \mathcal{V}\}| \quad (**)$$

We observe that adding $(**)$ over all components of H gives $\sum_{v \in H: S_v \in \mathcal{V}} |\delta_H(v)| \leq 2|\{v \in H : S_v \in \mathcal{V}\}|$.

Note that $\sum_{v \in C} |\delta_H(v)| = 2(|C| - 1)$ because it is a tree and then by handshaking lemma. So to show $(**)$, it suffices to show that if v is a leaf of C , then $S_v \in \mathcal{V}$. (so if $v \in C$, $S_v \notin \mathcal{V}$, then $|\delta_H(v)| \geq 2$.)

If C is a singleton component, $(**)$ clearly holds. Suppose C is not a singleton component, and let v be a leaf node node of C .



Here we also suppose $f(S_v) = 0$, i.e., $S_v \notin \mathcal{V}$. Since edge uv of C incident to v was not deleted, if $Z \subseteq V$ node-set corresponds to $C - \{v\}$ (i.e., $Z = \bigcup_{w \in C - \{v\}} S_w$), then $f(Z) = 1$.^a Then $f(V - Z) = 1$. But

$$V - Z = \left(V - \bigcup_{w \in C} S_w \right) \cup S_v$$

and $f(V - \bigcup_{w \in C} S_w) = f(\bigcup_{w \in C} S_w) = 0$ since C is a component of H . This contradicts property (iii) of proper functions.

To finish the proof,

$$\begin{aligned}
\sum_{v \in H: S_v \in \mathcal{V}} |\delta_H(v)| &= \sum_{\text{components } C \text{ of } H} \sum_{v \in C: S_v \in \mathcal{V}} |\delta_H(v)| \\
&= \sum_{\text{components } C \text{ of } H} \left(2|C| - 2 - \sum_{v \in C: S_v \notin \mathcal{V}} \underbrace{|\delta_H(v)|}_{\geq 2} \right) \\
&\leq \sum_{\text{components } C \text{ of } H} \left(2|C| - 2 - \sum_{v \in C: S_v \notin \mathcal{V}} 2 \right) \\
&= \sum_{\text{components } C \text{ of } H} (2|C| - 2|v \in C : S_v \notin \mathcal{V}| - 2) \\
&= \sum_{\text{components } C \text{ of } H} (2|v \in C : S_v \in \mathcal{V}| - 2) \\
&\leq \sum_{\text{components } C \text{ of } H} (2|v \in C : S_v \in \mathcal{V}|) \\
&= \sum_{\text{components } C \text{ of } H} \sum_{v \in C: S_v \in \mathcal{V}} 2 \\
&= 2|\{v \in H : S_v \in \mathcal{V}\}|
\end{aligned}$$

^aWe have $f(Z \cup S_v) = 0$ as C is a component of H . So there are no edges in final set F from $\delta(Z \cup S_v)$.

uv as not deleted because deleting uv will create an infeasible solution, which means some component S of $(V, F - uv)$ has $f(S) = 1$. As F is feasible, then such an S must be a new component that gets created when we delete uv : either $S = S_v$ or $S = Z$. As $f(S_v) = 0$, then $f(Z) = 1$.

Theorem 9.3

The final edge set F , dual solution y returned satisfy $c(F) \leq 2 \sum_{S \in \mathcal{S}} y_S$. Thus 2-approximation.

Proof:

Divide the algorithm into stages, where each stage i corresponds to an iteration of WHILE loop, and on associated time interval. Note that in a stage, collection of MVSs does not change. Let

Δ_i = length of stage i (i.e., length of associated time interval)

\mathcal{V}_i = collection of MVSs at start of stage i

Note that $y_S = \sum_{i: S \in \mathcal{V}_i} \Delta_i \quad \forall S \in \mathcal{S}$, which implies

$$\sum_{S \in \mathcal{S}} y_S = \sum_{S \in \mathcal{S}} \left(\sum_{i: S \in \mathcal{V}_i} \Delta_i \right) = \sum_i \Delta_i |\mathcal{V}_i|$$

Then we have

$$\begin{aligned}
c(F) &= \sum_{e \in F} c_e = \sum_{e \in F} \left(\sum_{S \in \mathcal{S}: e \in \delta(S)} y_S \right) \\
&= \sum_{S \in \mathcal{S}} y_S |\delta_F(S)| \\
&= \sum_{S \in \mathcal{S}} \left(\sum_{i: S \in \mathcal{V}_i} \Delta_i \right) |\delta_F(S)| \\
&= \sum_i \Delta_i \left(\sum_{S \in \mathcal{V}_i} |\delta_F(S)| \right) \\
&\leq \sum_i \Delta_i \cdot 2|\mathcal{V}_i| && \text{By Lemma 9.2} \\
&= 2 \sum_{S \in \mathcal{S}} y_S
\end{aligned}$$

Set Cover & Uncapacitated Facility Location

Recall Set Cover: Given a universe U of n elements, and a collection \mathcal{S} of subsets of U . Each set $S \in \mathcal{S}$ has a weight $w_S \geq 0$. **Goal:** Find a min-weight collection of sets from \mathcal{S} whose union is U , which is called a set cover.

Recall:

- Even when all $w_S = 1$, Set Cover is NP-hard.
- Even vertex cover (special case, where sets are equivalent to nodes of a graph, U is edges) is NP-hard.

10.1 LP-relaxation for SET COVER and its dual

As usual, S index sets in \mathcal{S} , e index elements in U .

$$\begin{aligned} \min \sum_S w_S x_S & \quad (\text{SC-P}) \\ \text{s.t. } \sum_{S: e \in S} x_S & \geq 1 \quad \forall e \in U \\ x & \geq 0 \end{aligned} \quad (1)$$

where x_S indicates if S is picked.

$$\begin{aligned} \max \sum_e y_e & \quad (\text{SC-D}) \\ \text{s.t. } \sum_{e \in S} y_e & \leq w_S \quad \forall S \in \mathcal{S} \\ y & \geq 0 \end{aligned} \quad (2)$$

A $\{0, 1\}$ -solution to (SC-P) is precisely a set cover.

CS conditions

$$x_S > 0 \implies \sum_{e \in S} y_e = w_S \quad \forall S \in \mathcal{S} \quad (3)$$

$$y_e > 0 \implies \sum_{S: e \in S} x_S = 1 \quad \forall e \in U \quad (4)$$

Let $B = \max_e |\{S \in \mathcal{S} : e \in S\}|$.

Claim For any feasible solution x to (SC-P), where $x_S \leq 1 \forall S$, we have $\sum_{S:e \in S} x_S \leq B$ for every $e \in U$. This is simply because there are at most B terms on LHS, each is at most 1.

So if we can construct a $\{0,1\}$ -solution \hat{x} feasible for (SC-P), and a dual feasible solution y such that (3) holds, then

$$\sum_S w_S \hat{x}_S = \sum_{S:\hat{x}_S=1} w_S = \sum_{S:\hat{x}_S=1} \left(\sum_{e \in S} y_e \right) = \sum_e y_e \left(\sum_{S:\hat{x}_S=1, e \in S} 1 \right) \leq B \sum_e y_e \leq B \cdot \text{OPT}(\text{SC-D})$$

which is B -approximation. This suggests following primal-dual B -approximation algorithm.

Algorithm 10: Primal Dual algorithm for Set Cover

```

1 Initially  $y \leftarrow 0, \mathcal{C} \leftarrow \emptyset, N = U - \sum_{S \in \mathcal{C}} S$  //  $\mathcal{C}$  is collection of sets picked.  $N$  is
   uncovered elements.
2 while  $N \neq \emptyset$  do
   Pick some  $e \in N$ 
   Increase  $y_e$  until (2) goes tight for some set  $S$  s.t.  $e \in S$ 
    $\mathcal{C} \leftarrow \mathcal{C} \cup \{S\}$ 
   Update  $N$ 
3 return  $\mathcal{C}$ 

```

Theorem 10.1

This primal dual algorithm is a B -approximation algorithm.

For vertex cover, $B = 2$ because each edge has 2 end points. Thus we have a 2-approximation for vertex cover.

Note that B can be very large, so the algorithm would be bad. Here is another algorithm for Set Cover.

Algorithm 11: Greedy algorithm for Set Cover

```

1  $y \leftarrow 0, \mathcal{C} \leftarrow \emptyset, N \leftarrow U$ . Initialize "time"  $t \leftarrow 0$  //  $t$  increase at rate 1 throughout
2 while  $N \neq \emptyset$  do
   Raise  $y_e \forall e \in N$  uniformly at rate 1, until some set  $S$  is paid by elements in  $N$ , i.e.,
    $\sum_{e \in N \cap S} y_e = w_S$ .
    $\mathcal{C} \leftarrow \mathcal{C} \cup \{S\}, N \leftarrow N - S$ .
3 return  $\mathcal{C}$ 

```

We make several observations

- At any time t , if N_t is the set of uncovered elements of time t , we have $y_e = t \forall e \in N_t$.
- So the *first* set to be picked is the set with smallest $w_S/|S|$ ratio.

In general, given that we have picked sets in \mathcal{C} , and set of uncovered elements N , the next set to be picked is the set $S \notin \mathcal{C}$ with smallest $\frac{w_S}{|S \cap N|}$ ratio, which is the spirit of greediness.

Now let's analyze the greedy approximation algorithm. By design, $\sum_{S \in \mathcal{C}} w_S = \sum_e y_e$. Does this suggest an optimal solution? But note that y need not to be feasible for (SC-D). We could have $\sum_{e \in S} y_e > w_S$ wince we ignore y_e 's for $e \in S - N$.

Let $\Delta = \max_{S \in \mathcal{S}} |S|$.

For an integer $k \geq 1$, we denote

$$H_k := 1 + \frac{1}{2} + \cdots + \frac{1}{k}$$

which is k -th Harmonic number, and the value is roughly $\log k$. We will show $y' = y/H_\Delta$ is a dual feasible solution, where y is the dual solution at end of greedy algorithm. Then this implies Greedy algorithm is a H_Δ -approximation algorithm:

$$\sum_{S \in \mathcal{S}} y_e = \sum_{e \in U} y_e = H_\Delta \cdot \left(\sum_{e \in U} y'_e \right) \leq H_\Delta \cdot \text{OPT}(\text{SC-D})$$

Lemma 10.2

$\sum_{e \in S} y_e \leq H_\Delta \cdot w_S$ for all $S \in \mathcal{S}$.

Proof:

Consider any set S . Order elements of S : e_1, \dots, e_k , where $k = |S|$, s.t. e_1 is covered before e_2 , e_2 before e_3 , and so on. How large could y_{e_1} be?

Let t_i be time when e_i is covered. So $y_{e_i} = t_i$ by design.

At time $t_1 - \epsilon$ ($\epsilon > 0$ small), all e_1, \dots, e_k are uncovered. So $|S|(t_1 - \epsilon) \leq w_S$, where LHS is total y -contribution from $|S \cap N|$, then $t_1 \leq \frac{w_S}{k} + \epsilon$. As this holds for $\epsilon > 0$, then $t_1 = y_{e_1} \leq w_S/|S|$.

Now consider y_{e_2} . At this time, all of e_2, \dots, e_k are uncovered. Then $(|S| - 1)(t_2 - \epsilon) \leq w_S$. This holds for all $\epsilon > 0$. Thus $y_{e_2} = t_2 \leq \frac{w_S}{|S| - 1}$.

Continuing this way, we get $y_{e_i} = t_i \leq \frac{w_S}{|S| - i + 1}$ for all $i = 1, \dots, k = |S|$. Thus

$$\sum_{e \in S} y_e = \sum_{i=1}^{|S|} y_{e_i} \leq \sum_{i=1}^{|S|} \frac{w_S}{|S| - i + 1} = w_S \left(\frac{1}{|S|} + \frac{1}{|S| - 1} + \cdots + 1 \right) = w_S \cdot H_{|S|} \leq w_S \cdot H_\Delta$$

As y/H_Δ is dual feasible solution, then

Theorem 10.3

Greedy algorithm is an H_Δ -approximation algorithm.

10.2 LP-Rounding Algorithms for Set Cover

Observation 1 If $x_S^* \geq \frac{1}{c}$, and we set $\hat{x}_S = 1$ in an integer solution then cost increase by a factor of at most c (relative to cost of x^*)

Observation 2 For any element e , there exists some set S containing e s.t. $x_S^* \geq \frac{1}{B}$, because there are at most B terms on the LHS of (1).

These two observations suggest LP-rounding B -approximation algorithm: Pick all sets S with $x_S^* \geq \frac{1}{B}$. In observation 2, we get a set cover. In observation 1, we get a solution of cost at most $B \cdot \text{OPT}(LP)$. Thus a B -approximation.

LP-rounding $O(\ln n)$ -approximation for set cover

where n is number of elements. Important technique: **Randomized Rounding**.

Idea Interpret $x_S^* \in [0, 1]$ as a probability, then x_S^* is the probability that set S is picked.

(A) Pick each set S independently with probability x_S^* .

Issues

1. We may not cover every element.
2. Can only say that expected cost of sets picked is bounded:

$$\text{Expected cost} = \sum_S w_S \cdot \Pr[S \text{ is picked}] = \sum_S w_S \cdot x_S^* = \text{OPT}(LP)$$

We will *only* deal with issue (1): Can we bound $\Pr[e \text{ is not bounded}]$ for some element e , and hence bound the probability “we do not have a set cover”. For an element $f \in U$,

$$\begin{aligned} \Pr[f \text{ is not covered}] &= \Pr[\text{no set } S \text{ containing } f \text{ is picked}] && \text{sets are picked independently} \\ &= \prod_{S: f \in S} (1 - x_S^*) && \text{as } 1 - x \leq e^{-x} \\ &\leq \prod_{S: f \in S} e^{-x_S^*} \\ &= e^{-\sum_{S: f \in S} x_S^*} && \sum_{S: f \in S} x_S^* \geq 1 \\ &\leq e^{-1} \end{aligned}$$

Suppose we repeat (A) for $2 \ln n$ rounds, each round is independent of all other rounds. Now

$$\begin{aligned} \Pr[f \text{ is not covered at the end of } 2 \ln n \text{ rounds}] &= \prod_{i=1}^{2 \ln n} \Pr[f \text{ is not covered in round } i] \\ &\leq \prod_{i=1}^{2 \ln n} e^{-1} \\ &= e^{-2 \ln n} \\ &= \frac{1}{n^2} \end{aligned}$$

Basic fact (union bound): $P(A \cup B) \leq P(A) + P(B)$. thus

$$\begin{aligned} \Pr[\text{we don't have a set cover after } 2 \ln n \text{ rounds}] &= \Pr[\exists f \in U \text{ s.t. } f \text{ is not covered after } 2 \ln n \text{ rounds}] \\ &\leq \sum_{f \in U} \Pr[f \text{ is not covered after } 2 \ln n \text{ rounds}] \\ &\leq n \cdot \frac{1}{n^2} \\ &= \frac{1}{n} \end{aligned}$$

Then the expected cost of sets picked:

$$E[\text{cost of sets picked after } 2 \ln n \text{ rounds}] \leq \sum_{i=1}^{2 \ln n} E[\text{cost of sets picked in round } i] = 2 \ln n \cdot \text{OPT}(LP)$$

So we have a randomized algorithm such that

1. Algorithm has a probability of error, i.e., we do not have a set cover.
2. Can only say that expected cost of solution is at most $2 \ln n \cdot \text{OPT}(LP)$.

Now we want to make error probability equal to zero. If after $2 \ln n$ rounds, sound element is not covered, then simply pick the min-weight set S_f containing f , and add that to our solution, i.e., the sets picked. Now

$$\begin{aligned} E[\text{cost of our solution}] &\leq 2 \ln n \cdot \text{OPT}(LP) + \sum_{f \in U} \Pr[f \text{ is not covered after } 2 \ln n \text{ rounds}] \cdot w_{S_f} \\ &\leq 2 \ln n \cdot \text{OPT}(LP) + n \cdot \frac{1}{n^2} \cdot \text{OPT}(LP) \\ &\leq \left(2 \ln n + \frac{1}{n}\right) \text{OPT}(LP) \end{aligned}$$

as $w_{S_f} \leq \text{OPT}(LP)$. This is because any feasible LP solution satisfies $\sum_{S: f \in S} x_S \geq 1$; then

$$\text{OPT}(LP) = \sum_{S: f \in S} x_S w_S \geq \sum_{S: f \in S} (\min_{S: f \in S} w_S) x_S = w_{S_f} \sum_{S: f \in S} x_S \geq w_{S_f}$$

10.3 Uncapacitated Facility Location

We are given a complete bipartite graph $G = (V = F \cup C, E)$ (so there is an edge $ij \forall i \in F, j \in C$). F is facilities, C is clients.

- Every facility i has an opening cost $f_i \geq 0$.
- For each facility $i \in F$, client $j \in C$, we have a connection/assignment cost $c_{ij} \geq 0$, which is the cost of assigning client j to facility i .

The goal is to choose a set $F' \subseteq F$ facilities to open, and assign each client $j \in C$ to an open facility $i(j) \in F'$ so as to minimize $\sum_{i \in F'} f_i + \sum_{j \in C} c_{i(j)j}$ which is the sum of *facility opening cost* + *client assignment/connection cost*.

Note that any open facility can serve any number of clients, this is uncapacitated.

Recall this is NP-complete. We then have *LP-relaxation*. i indexes facilities in F , j indexes clients in C . Variables:

- y_i indicates if facility i is open $\forall i \in F$.
- x_{ij} indicates if client j is assigned to facility $i \forall i \in F, j \in C$.

$$\min \sum_{i \in F} f_i y_i + \sum_{j \in C} \sum_{i \in F} c_{ij} x_{ij} \quad (\text{UFL-P})$$

$$\text{s.t. } \sum_i x_{ij} \geq 1 \quad \forall j \in C \quad (1)$$

$$x_{ij} \leq y_i \quad \forall i \in F, j \in C \quad (2)$$

$$x, y \geq 0$$

(1) means every j is assigned to a facility. (2) means j is assigned to an open facility.

$$\max \sum_{j \in C} \alpha_j \quad (\text{UFL-D})$$

$$\text{s.t. } \alpha_j - \beta_{ij} \leq c_{ij} \quad \forall i, j$$

$$\sum_{j \in C} \beta_{ij} \leq f_i \quad \forall i$$

$$\alpha, \beta \geq 0$$

UFL

Recall i indexes facilities in F , j indexes clients in C . Variables:

- y_i indicates if facility i is open $\forall i \in F$.
- x_{ij} indicates if client j is assigned to facility $i \forall i \in F, j \in C$.

$$\min \sum_{i \in F} f_i y_i + \sum_{j \in C} \sum_{i \in F} c_{ij} x_{ij} \quad (\text{UFL-P})$$

$$\text{s.t.} \quad \sum_i x_{ij} \geq 1 \quad \forall j \in C \quad (1)$$

$$x_{ij} \leq y_i \quad \forall i \in F, j \in C \quad (2)$$

$$x, y \geq 0$$

$$\max \sum_{j \in C} \alpha_j \quad (\text{UFL-D})$$

$$\text{s.t.} \quad \alpha_j - \beta_{ij} \leq c_{ij} \quad \forall i, j$$

$$\sum_{j \in C} \beta_{ij} \leq f_i \quad \forall i$$

$$\alpha, \beta \geq 0$$

11.1 Economic Interpretation of Dual LP

α_j is amount j is willing to pay to get itself assigned to a facility.

First, suppose all facilities are free, i.e., have 0 cost, or facility is not charging any payment from clients

(to produce service). Then, each client j will want to get assigned to facility i' such that $c_{i'j} = \min_{i \in F} c_{ij}$, i.e., pays an amount $\min_{i \in F} c_{ij}$. So we get a lower bound of $\sum_{j \in C} (\min_{i \in F} c_{ij})$ on OPT.

To get a better lower bound, suppose each facility i charges an amount $\beta_{ij} \geq 0$ (share of its opening cost charred to j) from client j (to provide service). These cost shares have to satisfy the fairness condition $\sum_{j \in C} \beta_{ij} \leq f_i \forall i \in F$.

Now, client j has to pay a *net cost* of $c_{ij} + \beta_{ij}$ to get assigned to facility i . Then client j would pay $\min_{i \in F} (c_{ij} + \beta_{ij})$ and so $\sum_{j \in C} \min_{i \in F} (c_{ij} + \beta_{ij})$ is a lower bound on OPT.

11.2 Metric UFL

Observation If all F_j^* -sets were pairwise disjoint (i.e., $F_j^* \cap F_k^* = \emptyset \forall j, k \in C, j \neq k$), then we can find F^* satisfies (*) of cost $\leq \sum_i f_i y_i^* \leq \text{OPT}(LP)$. By opening cheapest (in terms of f_i) facility in each F_j^* set.

Goal Pick a subset $C' \subseteq C$ of clients such that

1. F_j^* sets for all $j \in C'$ are disjoint.
2. any $k \in C - C'$ is “close” to a client in C'

Algorithm 12: Clustering algorithm to find C'

```

1  $L \leftarrow$  list of all clients sorted in increasing  $\alpha_j^*$ -order (break ties arbitrarily)
2  $C' \leftarrow \emptyset$ 
3 while  $L \neq \emptyset$  do
    Let  $j$  be first client in  $L$ .  $C' \leftarrow C' \cup \{j\}$ ,  $L \leftarrow L - \{j\}$ .
    for  $k \in L$  do
        if  $F_k^* \cap F_j^* \neq \emptyset$  then
             $L \leftarrow L - \{k\}$ , set  $\text{nbr}(k) = j$ 
4 return  $C'$ 

```

Algorithm 13: Metric UFL

```

1 Run clustering algorithm to find  $C'$ 
2  $F' \leftarrow \emptyset$ 
3 for  $j \in C'$  do
    Open  $i \in F_j^*$  with smallest  $f_i$ , i.e., add  $i$  to  $F'$ 
    Assigning  $j$  to  $i$ , and assign every  $k \in C - C'$  with  $\text{nbr}(k) = j$  to  $i$ 
4 return  $F'$  and corresponding client assignment

```

Now let's analyze the approximation guarantee.

Claim 11.1

If $k \in C - C'$, and $\text{nbr}(k) = j \in C'$, then $c_{jk} \leq \alpha_j^* + \alpha_k^*$, and $\alpha_j^* \leq \alpha_k^*$.

Proof:

Since \square

Lemma 11.2

- (a) If $j \in C'$, then it is assigned to a facility $i \in F'$ such that $c_{ij} \leq \alpha_j^*$.
- (b) If $k \in C - C'$, then it is assigned to $i \in F'$ such that $c_{ik} \leq 3\alpha_k^*$.

Proof:

For part (a), it's easy to see that since $j \in C'$, it is assigned to some $i \in F_j^*$, thus $c_{ij} \leq \alpha_j^*$. \square

Lemma 11.3

Cost of opening facilities is at most $\sum_i y_i^*$.

Proof:

The cost of cheapest facility i is

□

Theorem 11.4

Total cost is at most $4 \cdot \text{OPT}(LP)$.

which implies this is a 4-approximation algorithm.

Proof:

By Lemma 11.3, total facility-opening cost is at most $\sum f_i y_i^* \leq \text{OPT}(LP)$. By Lemma 11.2, total assignment cost is at most $3 \sum_j \alpha_j^* = 3 \cdot \text{OPT}(LP)$. So total cost is at most $4 \cdot \text{OPT}(LP)$. □